



HARARE INSTITUTE OF TECHNOLOGY
DEPARTMENT OF INFORMATION TECHNOLOGY
BTECH (HONS) INFORMATION TECHNOLOGY

**Satellite-Based Tobacco Crop Health Monitoring
System Using Machine Learning and NDVI Analysis**

BY
PRINCE RUPIYA
H220536M

HIT 400

SUPERVISOR: MR P. MANYONGANISE

This project report was submitted to the Harare Institute of Technology in
partial fulfilment of the Bachelor of Technology (Hons) Degree in
Information Technology.

2025

Declaration

I, Prince Rupiya (H220536M), declare that this project report is my own original work. It has not been submitted, in whole or in part, for any other degree or qualification at Harare Institute of Technology or elsewhere. Where I have drawn on the work of others, I have referenced and acknowledged it fully in accordance with the Harvard referencing system.

Student Signature: _____

Date: _____

Student Name: Prince Rupiya

Reg. Number: H220536M

Supervisor Signature: _____

Date: _____

Supervisor Name: Mr P. Manyonganise

Department: Information Technology

Dedication

*To the smallholder tobacco farmers of Zimbabwe,
whose resilience and dedication to the land inspired this work.*

*To my family, for their unwavering support and encouragement
throughout this academic journey.*

*To the future of Zimbabwean agriculture — may technology
serve the hands that feed our nation.*

Acknowledgements

My deepest thanks go to Mr P. Manyonganise, my supervisor, for his guidance, patience, and consistently sharp feedback. His knowledge of both information systems and agricultural technology shaped this project in ways I couldn't have managed alone.

I'm grateful to the Department of Information Technology at the Harare Institute of Technology for the academic foundation and the resources that made this research possible.

This work would not exist without the open-source tools it depends on — Next.js, React-Leaflet, scikit-learn, Google Earth Engine, and the Open-Meteo project. I owe a quiet debt to every developer who contributed to those projects.

To my colleagues and fellow students who sat through testing sessions, offered feedback, and kept my spirits up: thank you.

And to my family — your patience and sacrifice through this whole programme meant more than I can properly say.

Abstract

Tobacco is Zimbabwe’s most valuable cash crop, yet the smallholder farmers who grow it still rely on manual inspection — walking the field, looking for visible symptoms, and hoping problems haven’t already gone too far. By the time stress shows up as wilting or discolouration, yield loss is usually already happening. This project addresses that gap directly. It presents a satellite-based tobacco crop health monitoring system that pulls live Sentinel-2 multispectral imagery through Google Earth Engine, applies machine learning and temporal trend analysis, and delivers plain-language health assessments to farmers through an ordinary web browser.

At the core is a logistic regression classifier trained on satellite and weather observations. It sorts tobacco field health into three categories — *Healthy*, *Moderate Stress*, and *High Stress* — with 97% accuracy on held-out test data. Five features drive the model: mean NDVI, the trend slope of NDVI over time, within-field NDVI variance, average temperature, and cumulative rainfall across a 14-to-30-day window. The temporal trend component is particularly important: declining NDVI trajectories can be flagged days before visible symptoms emerge, giving farmers a real window to intervene. A within-field heatmap adds spatial detail, highlighting which parts of a field are under most stress.

The application runs on Next.js 14 with TypeScript. React-Leaflet handles the interactive map, and Prisma with PostgreSQL manages persistent storage. A rule-based recommendation engine sits on top of the classifier, converting raw analysis outputs into specific, readable advice on irrigation, disease risk, and field inspection.

Everything in this system — the satellite imagery, the weather data, the hosting platform — is free and open. There are no subscription fees. That was a deliberate design constraint, because the commercial precision agriculture tools that could theoretically serve Zimbabwean smallholder farmers are priced well beyond what those farmers can afford. This project shows that high-quality crop health monitoring doesn’t require expensive infrastructure. It requires thoughtful engineering.

Keywords: crop health monitoring, machine learning, NDVI, precision agriculture, satellite imagery, tobacco farming, Zimbabwe, Google Earth Engine, Sentinel-2, time-series analysis, web application.

Contents

Declaration	ii
Dedication	iii
Acknowledgements	iv
Abstract	v
List of Acronyms	xi
1 Introduction	1
1.1 Introduction	1
1.2 Motivation	2
1.3 Problem Statement	2
1.4 Related Work	3
1.4.1 Climate FieldView	3
1.4.2 CropIn Technology	3
1.4.3 Taranis	4
1.4.4 Google Earth Engine and Open Research Tools	4
1.5 Hypothesis	4
1.6 Technical Objectives	4
1.7 Expected Results	5
1.8 Ethics Considerations	6
1.9 Conclusion	6
2 Literature Review	8
2.1 Introduction	8
2.2 Related Works	8
2.2.1 FarmLogs (TELUS Agriculture)	8
2.2.2 SatAgro	9
2.2.3 Farmerline	9
2.2.4 Digital Green	9
2.2.5 Regrow Ag	10
2.2.6 Academic Foundations	10
2.3 Conclusion	11
3 Requirements Analysis	12

3.1	Introduction	12
3.1.1	Evaluate Alternatives	12
3.1.2	Outsourcing	12
3.1.3	Improvement Over the Current System	13
3.1.4	Development Approach	13
3.2	User Requirements	14
3.2.1	Data Collection Phase	15
3.2.2	Technical Feasibility	17
3.2.3	Hardware Requirements	17
3.2.4	Software Requirements	17
3.2.5	Technical Expertise	18
3.3	Economic Feasibility	18
3.3.1	Cost-Benefit Analysis	18
3.3.2	Tangible Benefits	19
3.3.3	Intangible Benefits	20
3.4	Operational Feasibility	20
3.4.1	Schedule Feasibility	20
3.5	Work Plan	20
3.5.1	Work Schedule	20
3.5.2	Gantt Chart	21
3.6	Conclusion	21
4	System Analysis	23
4.1	Introduction	23
4.2	Description of the Current System	23
4.3	Analysis of the Existing System	23
4.3.1	Context Diagram of the Existing System	24
4.3.2	Weaknesses of the Current System	24
4.4	Description of the Proposed Solution	25
4.4.1	Context Diagram and Data Flow Diagrams	25
4.5	Requirements Analysis	26
4.5.1	Functional Requirements	26
4.5.2	Non-Functional Requirements	28
4.6	System Models	29
4.6.1	UML Activity Diagram	29
4.6.2	UML Class Diagram	30
4.6.3	UML Sequence Diagram	32
4.7	Conclusion	32

5	System Design	34
5.1	Introduction	34
5.2	System Design	34
5.2.1	How will the System Work?	36
5.3	Solution Architecture	37
5.4	Database Modelling	38
5.4.1	E-R Diagram	38
5.4.2	Data Dictionary	39
5.4.3	Database Schema	41
5.5	Algorithm Design	42
5.5.1	Inputs to the Algorithm	42
5.5.2	Pre-processing and Validation	43
5.5.3	Feature Engineering	43
5.5.4	Model Selection and Responsibilities	44
5.5.5	Inference and Post-processing Logic	44
5.5.6	Algorithm Pseudocode	45
5.6	Interface Design	46
5.6.1	Login and Registration Interface	46
5.6.2	Dashboard Interface	47
5.6.3	Create Field Interface	47
5.6.4	Analysis Results Interface	48
5.7	Conclusion	49
6	System Implementation	50
6.1	Introduction	50
6.2	Development Environment and Tools	50
6.3	Application Structure	52
6.4	Frontend Implementation	52
6.4.1	Interactive Map Component	52
6.4.2	Analysis Results Page	53
6.5	Backend API Implementation	53
6.5.1	Authentication Routes	53
6.5.2	Field Management Routes	54
6.5.3	Analysis Route	54
6.6	Google Earth Engine Integration	55
6.6.1	Authentication and Service Account	55
6.6.2	NDVI Retrieval Pipeline	55
6.6.3	Heatmap Data Generation	55
6.7	Machine Learning Implementation	56

6.7.1	Training Pipeline	56
6.7.2	Model Export and JavaScript Inference	56
6.7.3	Recommendation Engine	57
6.8	Database Implementation	57
6.9	Conclusion	57
7	Testing, Evaluation and Conclusion	58
7.1	Introduction	58
7.2	Summary of Project Achievements	58
7.3	Evaluation of Technical Objectives	59
7.4	System Validation and Testing	60
7.4.1	Machine Learning Model Performance	60
7.4.2	Recommendation Engine Validation	61
7.4.3	Functional Testing	62
7.4.4	Non-Functional Testing	62
7.5	Limitations of the Current Implementation	63
7.6	Recommendations for Future Work	64
7.7	Conclusion	65
	References	66
A	Python Model Training Script	69
B	Model Export Script	72
C	JavaScript Inference Module	74
D	Recommendation Engine Source	76
E	Prisma Database Schema	78
F	Threshold and Model Configuration Files	81

List of Figures

3.1	Project Gantt Chart (March 2024 – February 2025)	21
4.1	Context Diagram — Existing Tobacco Crop Monitoring System	24
4.2	Context Diagram — Proposed Satellite Crop Health Monitoring System	26

4.3	Level 1 Data Flow Diagram — Proposed System	26
4.4	Use Case Diagram — Satellite Crop Health Monitoring System	28
4.5	UML Activity Diagram — Crop Health Analysis Workflow	30
4.6	UML Class Diagram — System Data Model	31
4.7	UML Sequence Diagram — Crop Health Analysis Request	32
5.1	Proposed System Architecture — TobaccoGuard	37
5.2	Entity-Relationship Diagram — TobaccoGuard Database	39
5.3	Wireframe — Login Page	46
5.4	Wireframe — Dashboard (Field List View)	47
5.5	Wireframe — Create Field (Interactive Map View)	48
5.6	Wireframe — Analysis Results Page	49

List of Tables

3.1	Two-Year Cost-Benefit Analysis (USD)	19
4.1	Functional Requirements	27
6.1	Development Tools and Technologies	51
6.2	Application Directory Structure	52
7.1	Evaluation of Technical Objectives Against Implementation	59
7.2	NDVI and Weather Classification Thresholds	61
7.3	Classification Report — Logistic Regression Model on Test Set ($n = 122$)	61
7.4	Recommendation Engine Rules and Trigger Conditions	62
A.1	Feature Configuration — <code>config.py</code>	71

List of Acronyms

Acronym	Definition
API	Application Programming Interface
CRIM	Crop Remote Intelligence Monitor
DFD	Data Flow Diagram
ESA	European Space Agency
EVI	Enhanced Vegetation Index
FAO	Food and Agriculture Organization of the United Nations
GEE	Google Earth Engine
GIS	Geographic Information System
HTTP	Hypertext Transfer Protocol
HIT	Harare Institute of Technology
IT	Information Technology
JSON	JavaScript Object Notation
L2A	Level-2A (surface reflectance product of Sentinel-2)
ML	Machine Learning
NDVI	Normalized Difference Vegetation Index
NIR	Near-Infrared
ORM	Object-Relational Mapper
PostGIS	PostgreSQL Geographic Information System Extension
REST	Representational State Transfer
SQL	Structured Query Language
TIMB	Tobacco Industry and Marketing Board
UML	Unified Modelling Language
UI	User Interface
URL	Uniform Resource Locator
USD	United States Dollar

Chapter 1: Introduction

1.1 Introduction

Agriculture is the backbone of Zimbabwe's economy. It accounts for roughly 12–15% of GDP and keeps an estimated 60–70% of the rural workforce employed ([Ministry of Agriculture, 2022](#)). Within that sector, tobacco is the crown jewel — the country's most valuable cash crop and one of the largest tobacco outputs anywhere in the world. The Tobacco Industry and Marketing Board recorded auction sales exceeding 184 million kilogrammes in the 2022/2023 season alone, generating over USD 600 million in export revenue ([TIMB, 2023](#)). For smallholder farmers on communal and resettlement lands, a good tobacco season is the difference between financial stability and debt.

Yet despite that strategic importance, most of those farmers are still monitoring their crops the same way they always have: walking the field, looking for symptoms, and relying on experience to guide decisions ([Nyoka et al., 2020](#)). It's a method that works — until it doesn't. Manual inspection is slow, subjective, and limited by how much of the field a person can realistically cover. The deeper problem is that it only catches stress after damage has already become visible. By the time wilting or discolouration shows up, the plant has already been under physiological stress for days. That translates directly to leaf quality loss, lower auction grades, and less money in the farmer's pocket ([Rwizi et al., 2021](#)).

Satellite remote sensing changes that equation. The Normalized Difference Vegetation Index (NDVI), derived from multispectral imagery, has been used to track vegetation health for over four decades ([Rouse et al., 1974](#)). It's computed from the ratio of near-infrared to red reflectance and is sensitive to changes in plant physiology that occur well before any visible symptoms appear — often by several days. Paired with weather data and a machine learning classifier, NDVI time-series become an early-warning system rather than just a record.

This project builds exactly that. It implements a Satellite-Based Tobacco Crop Health Monitoring System that pulls live Sentinel-2 multispectral imagery through Google Earth Engine, runs it through a logistic regression classifier and temporal trend analysis, and presents the results through a web interface any farmer can use on a smartphone. Draw a field boundary on a map, trigger the analysis, and get specific recommendations on irrigation, disease monitoring, and crop management — all driven by actual satellite data and at zero cost.

The chapters that follow build out the full picture. Chapter 1 establishes the context, motivation, and objectives. Chapter 2 surveys the existing landscape of precision agriculture

platforms and the academic foundations this project draws on. Chapter 3 covers requirements and feasibility. Chapter 4 presents the full system analysis, from existing-system critique through to formal UML models.

1.2 Motivation

Two realities collide here, and together they make this project worth doing.

The first is economic. Tobacco is Zimbabwe’s single biggest foreign currency earner in agriculture ([TIMB, 2023](#)), yet yield variability among smallholder growers remains stubbornly high. A significant part of that variability comes down to late detection — stress events that could have been caught early instead go unnoticed until they’ve already cost the farmer money. The Ministry of Lands, Agriculture, Fisheries, Water and Rural Development has recognised this, naming precision farming as a strategic priority in the Zimbabwe Agricultural Development Strategy II 2021–2025 ([Ministry of Agriculture, 2022](#)).

The second reality is technological. Satellite data that used to require expensive hardware and specialised processing infrastructure is now freely available through Google Earth Engine ([Gorelick et al., 2017](#)). ESA’s Sentinel-2 satellites cover Zimbabwe every 5–10 days at 10-metre resolution ([ESA, 2023](#)) — more than enough to catch a developing stress event. The data exists and it’s free. What’s missing is an interface that makes it usable by a farmer without a remote sensing background.

That’s the gap this project fills. Translating freely available satellite data into actionable farm-level insights — without requiring specialist training or expensive equipment — is both technically feasible and genuinely needed. Detecting a declining NDVI trend 14 days before symptoms appear gives farmers a window to intervene. Manual inspection can’t do that.

1.3 Problem Statement

The problems with current monitoring practice fall into six distinct categories:

1. **Reactive detection:** Stress — from water deficit, pest infestation, or nutrient deficiency — only gets noticed once symptoms are visible. At that point, yield loss is already locked in ([Nyoka et al., 2020](#)).
2. **Temporal gaps:** Inspections happen weekly or fortnightly at best. A stress event that starts on day 3 and intensifies by day 10 can cause irreversible damage before anyone notices.
3. **Spatial gaps:** Even a thorough field visit covers only a fraction of the actual field area. A waterlogged corner, a patch of compacted soil, or a pest entry point at the far

edge of a field can go entirely undetected.

4. **No environmental context:** Monitoring decisions are made without looking at weather data. Rainfall totals, temperature anomalies, and drought indices that strongly predict stress events simply aren't factored in.
5. **Cost barriers:** Commercial satellite monitoring platforms charge hundreds to thousands of USD per year ([Wolfert et al., 2017](#)). That pricing puts them out of reach for smallholder farmers whose entire annual income from a season may be comparable.
6. **No local calibration:** Very few systems have been adapted for Zimbabwean growing conditions, tobacco-specific NDVI thresholds, or the practical realities of smallholder farming in Southern Africa ([Rwizi et al., 2021](#)).

Taken together, these gaps make the case for a system that monitors tobacco crop health continuously from satellite data, flags stress conditions before they become visible, and communicates findings in terms farmers can act on.

1.4 Related Work

A number of platforms and research tools are relevant here, both as prior art and as context for what's still missing.

1.4.1 Climate FieldView

Climate FieldView — developed by The Climate Corporation, a Bayer AG subsidiary — is among the most widely deployed commercial precision agriculture platforms in the world ([Climate Corporation, 2023](#)). It covers satellite field imagery, NDVI monitoring, yield mapping, and farm machinery integration. The technical capability is not in question. But FieldView is built for large-scale North American grain operations, its NDVI thresholds and recommendation logic are calibrated for maize and soybean, and its subscription costs make it economically inaccessible for Zimbabwean smallholders. None of that is fixable with configuration.

1.4.2 CropIn Technology

CropIn, an Indian agri-tech firm, has deployed digital farming tools in parts of sub-Saharan Africa through NGO and government partnerships ([CropIn Technology, 2023](#)). Its coverage of crop management, weather monitoring, and field analytics shows that digital agriculture can work in developing-region contexts. What CropIn doesn't offer is tobacco-specific calibration, validated performance against Zimbabwean field conditions, or access without paid subscription tiers. It's a useful proof-of-concept for the approach, not a ready solution.

1.4.3 Taranis

Taranis uses high-resolution drone and satellite imagery to detect crop disease and pest damage at near-leaf-level precision ([Taranis, 2023](#)). The technical depth is impressive — it detects things standard satellite imagery can't. But it's designed for large commercial farms, requires drone hardware on the ground, and prices per hectare in a way that simply doesn't work for smallholder contexts. Taranis is a good example of a technically strong tool with the wrong access model.

1.4.4 Google Earth Engine and Open Research Tools

GEE itself, along with tools like the Open Data Cube and Sen2Agri, provides the kind of free satellite processing infrastructure that serious agricultural remote sensing research runs on ([Gorelick et al., 2017](#)). [Rwizi, Sibanda and Chimonyo \(2021\)](#) have demonstrated GEE-based NDVI monitoring working in Zimbabwean agricultural contexts, confirming that the technical approach is viable locally. The limitation is that these tools live in the research domain — they require programming knowledge and offer no farmer-facing interface.

1.5 Hypothesis

Given a web-based system that combines live Sentinel-2 NDVI data retrieved via Google Earth Engine with machine learning classification and temporal trend analysis, Zimbabwean smallholder tobacco farmers will be able to detect crop stress conditions 7 to 14 days before symptoms become visible. That early warning window, if acted upon, should reduce crop losses, improve how inputs are allocated, and increase both yield and income.

1.6 Technical Objectives

The overarching aim is to design and implement a Satellite-Based Tobacco Crop Health Monitoring System that smallholder farmers in Zimbabwe can actually use — accessible, accurate, and deployable without specialist infrastructure. The specific objectives are:

1. Build a web platform that lets farmers draw field polygons on an interactive map and trigger a satellite-based crop health analysis on demand.
2. Pull live Sentinel-2 Level-2A imagery through the Google Earth Engine API to compute real NDVI values for each user-defined polygon, replacing estimated figures with actual satellite measurements.

3. Train a logistic regression classifier on ground-truth satellite and weather observations to categorise tobacco field health as *Healthy*, *Moderate Stress*, or *High Stress*.
4. Implement temporal NDVI trend analysis across a 14-to-30-day window, so that declining vegetation trajectories can be flagged before fields reach a critical stress state.
5. Compute within-field spatial variability from per-pixel NDVI distributions and render it as a heatmap overlay, enabling farmers to pinpoint exactly where stress is concentrated.
6. Build a rule-based recommendation engine that converts the classifier output and raw feature values into plain-language advice on irrigation, disease monitoring, and field inspection.
7. Architect the system so Python handles model training offline and JavaScript handles inference at runtime — making the whole thing deployable serverlessly, without a Python backend.

1.7 Expected Results

The expected outcomes of a successful implementation are:

- A working web application — no installation required — that gives Zimbabwean tobacco farmers access to live satellite-based crop health monitoring from any browser.
- A logistic regression classifier clearing 90% accuracy on held-out test data, with balanced precision and recall across all three health classes.
- Trend detection that catches declining NDVI trajectories within a 14-to-30-day window, before fields cross into the High Stress category.
- A within-field heatmap showing exactly where stress is concentrated inside a drawn polygon, so interventions can be targeted rather than applied uniformly.
- A recommendation engine that links each piece of advice — whether on irrigation, disease risk, or calling in an agronomist — to the specific data reading that triggered it.
- A zero-cost data architecture built entirely on GEE’s free research tier and the Open-Meteo API, with no recurring subscription fees.
- Better-informed farm decisions, with the reasonable expectation that catching stress 7–14 days early translates into fewer crop losses and better use of inputs.

1.8 Ethics Considerations

- **Data Privacy and Consent:** The system collects names, email addresses, and field location data. All of it is stored in a PostgreSQL database protected by role-based access controls and bcrypt-hashed passwords. Farmers own their field data; nothing is shared with third parties without explicit consent. The system operates in accordance with Zimbabwe’s Cybersecurity and Data Protection Act (2021).
- **Data Accuracy and Transparency:** Every analysis result comes with its observation date, cloud cover status, and analysis window parameters so users can judge data quality themselves. Where cloud cover has reduced the number of usable satellite scenes, the system says so explicitly. Each recommendation is tied back to the data reading that produced it, so farmers can evaluate the advice rather than just accept it.
- **Algorithmic Fairness:** The model is trained on data that reflects real variation in Zimbabwean tobacco field conditions. Performance is evaluated class by class — not just overall accuracy — to catch any category-specific bias early. The NDVI thresholds are cross-checked against agronomic literature to confirm they correspond to genuine physiological boundaries, not just statistical artefacts.
- **Accessibility and Inclusivity:** The interface is built for basic internet connectivity and entry-level smartphones. Text is written in plain English, pitched at secondary school reading level. No technical literacy is assumed beyond the ability to navigate a standard website.
- **Environmental Responsibility:** By enabling targeted irrigation and agrochemical application to specific stressed zones rather than blanket treatment of the whole field, the system actively reduces water and chemical waste.
- **Satellite Data Use:** All imagery is accessed through GEE under its terms of service for research and non-commercial use. Sentinel-2 data is released freely by ESA under the Copernicus Open Access Data Policy with no use restrictions.

1.9 Conclusion

The challenges are real and well-documented: Zimbabwean smallholder tobacco farmers detect stress too late, inspect too little of their fields, work without environmental data, and face subscription costs that make commercial solutions irrelevant to them. None of those problems are unsolvable. Freely available Sentinel-2 satellite imagery, machine learning classification, and a sensibly designed web interface are sufficient to address all of them.

What follows from here is the substance: a literature review that grounds the approach in what's already been tried; a requirements analysis that connects the system design to actual farmer needs; and a system analysis that works through the architecture in formal detail.

Chapter 2: Literature Review

2.1 Introduction

Precision agriculture has come a long way since the early NDVI monitoring work of the 1970s. What started as research-grade tools that only remote sensing specialists could use has grown into a mature industry: integrated platforms that combine satellite imagery, field sensors, weather data, and machine learning in commercially deployed products ([Wolfert et al., 2017](#)). The question for this project isn't whether the technology works — it does — but whether any of it works for smallholder tobacco farmers in Zimbabwe.

This chapter works through five relevant platforms — FarmLogs, SatAgro, Farmerline, Digital Green, and Regrow Ag — assessing each for technical capability, accessibility, and fit with the Zimbabwean smallholder context. It also reviews the academic literature on NDVI monitoring, machine learning for agriculture, and decision support system design that underpins the methodology used in this project.

2.2 Related Works

2.2.1 FarmLogs (TELUS Agriculture)

FarmLogs — now part of TELUS Agriculture — is a cloud-based farm management platform covering satellite field imagery, NDVI crop health monitoring, and weather-based planting recommendations ([Farmerline, 2023](#)). It integrates with GPS-tracked equipment and soil sampling records, and it presents NDVI as a time-series health index derived from Landsat and Sentinel-2 data. Technically, it's a solid demonstration that web-based NDVI monitoring works.

The problem is everything else. Subscriptions run USD 300 to USD 600 per year ([Wolfert et al., 2017](#)) — in the same ballpark as a Zimbabwean smallholder farmer's entire net income from a season. Beyond cost, the platform is built for North American grain farming. Its thresholds, recommendation logic, and calibration are designed for maize and soybean in temperate climates, not tobacco in Zimbabwe's tobacco belt. [Liakos et al. \(2018\)](#) make the point plainly: an agricultural decision support system that hasn't been calibrated for the crop and context it's being applied to simply isn't useful. FarmLogs hasn't been.

2.2.2 SatAgro

SatAgro is a European service that uses Sentinel-1 and Sentinel-2 imagery to produce NDVI time-series, field zone maps, and variable-rate fertiliser prescription maps for cereal and oilseed farming ([SatAgro, 2023](#)). It computes per-polygon NDVI statistics and generates within-field variability maps — exactly the kind of spatial hotspot analysis this project implements.

That makes SatAgro the closest commercial analogue to one part of what this system does. But the similarity ends there. SatAgro operates in European markets only, is calibrated for temperate-climate crops, and has no provision for the cloud-cover gaps that characterise Zimbabwe’s rainy season from November to March. It’s also subscription-based. Nonetheless, the fact that within-field variability mapping is a commercially proven capability in another context validates building it here.

2.2.3 Farmerline

Farmerline is a Ghanaian company that has deployed digital farming tools across West and East Africa, reaching smallholder farmers through smartphone apps and interactive voice response systems for feature phones ([Farmerline, 2023](#)). It provides weather forecasts, crop management advice, and market price information — and it actually reaches the kinds of farmers this project is designed for. In that sense it’s more contextually relevant than anything in the commercial platform category above.

What Farmerline doesn’t do is satellite monitoring. Its advice comes from crop calendars and generalised weather data, not from field-specific observations. [Rose et al. \(2016\)](#) note that this is a fundamental limitation: decision support tools built on generic rules without field-level data can’t be specific enough to be genuinely useful on an individual farm. That said, Farmerline’s success at getting farmers to actually adopt the tool is instructive. Plain language, minimal jargon, and a design that works on entry-level hardware — those are the interface principles this project takes directly from Farmerline’s playbook.

2.2.4 Digital Green

Digital Green is an NGO-driven platform that uses video-based peer learning and mobile advisory tools to improve smallholder farming practice across South Asia and Sub-Saharan Africa, including Zimbabwe ([Digital Green, 2023](#)). It has worked directly with Zimbabwe’s Ministry of Agriculture to spread good agricultural practice through community video screenings. The yield improvements it has demonstrated in smallholder communities are real evidence that technology-mediated extension works in the Zimbabwean context.

But Digital Green doesn’t do satellite monitoring, NDVI analysis, or field-level health

assessment. Its role is disseminating general best practices, not providing field-specific precision data. The two functions aren't in competition — a farmer who uses Digital Green for agronomic education and this system for field-level health monitoring is better served than one using either alone.

2.2.5 Regrow Ag

Regrow Ag uses satellite data, crop simulation models, and machine learning to measure sustainable crop production and quantify carbon sequestration for large agricultural supply chains (Regrow Ag, 2023). The processing pipeline is technically sophisticated — it handles Sentinel-2 and Landsat imagery at scale. But the outputs go to corporate sustainability dashboards, not to farmers. Regrow Ag is built for commodity companies and cooperatives, not for individuals.

The interesting thing about Regrow Ag, in the context of this project, is that it's the only platform reviewed here that explicitly incorporates temporal NDVI trend analysis. It computes seasonal vegetation trajectories to estimate crop productivity. That's validation for the core analytical approach this project uses. The gap Regrow Ag leaves open is a farmer-facing version of that same capability — which is precisely what this project fills.

2.2.6 Academic Foundations

The academic literature provides the methodological backbone for this project's approach.

Rouse et al. (1974) established NDVI as a reliable vegetation health metric by demonstrating its correlation with biomass and photosynthetic activity. Tucker (1979) then confirmed the physiological basis: NDVI responds to chlorophyll absorption in the red band and canopy scattering in the near-infrared in ways that precede visible stress symptoms — making it genuinely useful as an early-warning indicator rather than just a diagnostic one.

Weiss et al. (2020) meta-reviewed 110 remote sensing studies in agriculture and found that satellite-derived vegetation indices reliably classify crop health status, with prediction accuracy exceeding 80% in most controlled settings. They identify Sentinel-2 — 10-metre resolution, 5-to-10-day revisit — as the best freely available data source for field-scale monitoring, which is why this project uses it.

On machine learning, Liakos et al. (2018) reviewed 40 agriculture studies and found that logistic regression, SVMs, and random forests dominate the literature. Critically, they note that logistic regression — despite its apparent simplicity — frequently matches more complex models when the features are well-engineered and class boundaries are reasonably stable. That observation directly informs the algorithm choice here. Chlingaryan et al. (2018) add that combining NDVI with weather variables (temperature and precipitation) consistently outperforms NDVI-only models for crop health classification, which is why the feature set in this project includes both.

Rose et al. (2016) reviewed 40 agricultural decision support tools and found that farmer adoption correlates strongly with transparency: tools that explain their reasoning get trusted and used; black-box tools don't. That finding shapes the entire recommendation engine in this project. And Wolfert et al. (2017), reviewing big data architectures in precision agriculture, make the case for separating data acquisition, analysis, and presentation — the three-layer design principle implemented here.

2.3 Conclusion

The pattern across the commercial landscape is consistent: the platforms with the strongest technical capability — FarmLogs, SatAgro, Regrow Ag — are built for large operations in developed markets and priced accordingly. The platforms that have genuinely reached smallholder farmers in developing regions — Farmerline, Digital Green — don't do satellite monitoring or field-level precision analysis. There's no platform that does both.

More specifically, nothing in this review combines live satellite NDVI analysis, temporal trend detection, within-field spatial variability mapping, and crop-specific recommendation logic in a single free, farmer-facing web interface calibrated for Zimbabwean conditions.

The academic literature validates the individual technical components: Sentinel-2 NDVI is a proven early-warning indicator; logistic regression on NDVI and weather features works; explainable systems get adopted at higher rates than opaque ones; and temporal trend analysis is genuinely valuable but largely absent from farmer-facing tools. Together, these findings define exactly what needs to be built.

Chapter 3: Requirements Analysis

3.1 Introduction

Getting requirements wrong at the start is expensive. [Sommerville \(2016\)](#) makes the point that errors introduced during requirements engineering cost more with every subsequent phase that builds on top of them. For this project, requirements were gathered through direct observation of current tobacco monitoring practice, structured interviews with farmers and extension officers, a questionnaire distributed to 20 stakeholders, analysis of existing systems and published literature, and early prototyping exercises with prospective users.

This chapter evaluates the development alternatives that were considered before custom-build was chosen, defines the requirements for three user roles, assesses technical, economic, and operational feasibility, and closes with the project work plan.

3.1.1 Evaluate Alternatives

Two alternatives were examined before committing to a custom build.

The first was to configure an existing commercial precision agriculture platform. The platforms reviewed in Chapter 2 — FarmLogs and SatAgro in particular — were assessed against the Zimbabwean smallholder tobacco context and rejected on three grounds: subscription pricing that farmers can't sustain; no tobacco-specific NDVI thresholds or recommendation logic; and no deployment or support infrastructure in Zimbabwe. [Wolfert et al. \(2017\)](#) note that commercial platforms built for developed-market farming regularly fail in smallholder technology transfer precisely because of these structural mismatches.

The second was to build an extension layer on an open-source GIS platform like QGIS or Open Data Cube. That would have given access to satellite processing infrastructure, but would have required users to install and configure specialised GIS software — completely at odds with the accessibility requirement for farmers with basic digital literacy. Neither option offered a machine learning inference pipeline, a recommendation engine, or a deployable web interface.

Custom development was the right choice, and that's what was built.

3.1.2 Outsourcing

Outsourcing was also considered and ruled out. The system requires tobacco-specific agonomic knowledge, close familiarity with the GEE API for Zimbabwean satellite data, and the kind of iterative refinement that only works when the developer is deeply embedded in

the problem. [Pfleeeger and Atlee \(2010\)](#) note that outsourcing works best when requirements are fully specified and stable upfront — which isn’t the case in a research and prototyping context where requirements evolve through building and testing. Keeping the codebase in-house is also a prerequisite for the academic contribution and any future open-source release.

3.1.3 Improvement Over the Current System

Five specific improvements over manual monitoring:

- **Timeliness:** Sentinel-2 revisit cycles give satellite analysis every 5–10 days, versus the weekly or fortnightly inspections that pass for “frequent” monitoring under current practice.
- **Objectivity:** NDVI values come from satellite reflectance measurements. They’re quantitative and reproducible, not dependent on how tired the inspector is that day.
- **Early detection:** Temporal trend analysis flags declining NDVI trajectories 7–14 days before anything shows up visually, giving farmers an actual intervention window.
- **Spatial precision:** Within-field heatmapping shows exactly where stress is concentrated — something a walking inspection covering three to five points in a field reliably misses.
- **Cost:** No recurring fees. The system uses exclusively free data sources, and automated monitoring replaces at least some of the physical inspection visits.

[Pfleeeger and Atlee \(2010\)](#) make the point that improvements which directly affect decision timeliness and accuracy aren’t cosmetic — they’re core requirements. That principle runs through every design decision in this project.

3.1.4 Development Approach

The system uses a three-layer architecture: data acquisition, machine learning analysis, and user interface. Each layer can be developed, tested, and updated independently. The ML model trains offline in Python with scikit-learn ([Pedregosa et al., 2011](#)) and gets exported as a JSON weight file, so the Next.js application can run inference in JavaScript without needing a Python server at all. The web layer is Next.js 14 with the App Router, TypeScript throughout, and Prisma as the ORM. Development follows an Agile approach — two-week sprints, iterative prototyping, weekly supervisor check-ins.

3.2 User Requirements

Three user roles interact with the system:

Farmer Requirements:

- Register and log in with email and password.
- Draw field boundary polygons on an interactive satellite map.
- Trigger a crop health analysis for any saved field with a single action.
- Receive analysis results in plain language with colour-coded health status indicators.
- Get specific, actionable recommendations covering irrigation, disease monitoring, and field inspection.
- View a history of previous analyses for each field.
- Access all of the above on a smartphone browser with mobile data.

Agronomist Requirements:

- View satellite analysis results for assigned fields.
- See NDVI time-series trend data in chart form.
- Access within-field heatmap visualisations of NDVI spatial variability.
- Inspect the raw feature values (mean NDVI, trend slope, variance, temperature, rainfall) behind each classification.
- Browse analysis history across multiple growing seasons.

System Administrator Requirements:

- Manage user accounts: create, assign roles, deactivate.
- Configure system-level parameters including NDVI threshold values and recommendation rule weights.
- Access usage logs for auditing and monitoring.
- Enforce secure authentication and role-based access controls across all administrative functions.

3.2.1 Data Collection Phase

Requirements were gathered through four methods (Sommerville, 2016):

Structured Interviews: Semi-structured interviews with five tobacco farmers in the Mvurwi area of Mashonaland Central Province and three Agritex extension officers. The conversations covered current monitoring practice, inspection frequency, awareness of satellite tools, and appetite for a digital alternative. All five farmers said they detect water stress only once wilting is already visible. None had heard of NDVI. All wanted a system that could warn them earlier.

Questionnaire: A structured questionnaire distributed to 20 stakeholders — smallholder farmers, cooperative managers, and extension officers. It is reproduced below.

Sample Questionnaire: Tobacco Crop Monitoring Needs Assessment

Purpose: This questionnaire gathers information on current crop monitoring practices and technology readiness among tobacco farming stakeholders in Zimbabwe, in support of the development of a Satellite-Based Tobacco Crop Health Monitoring System. All responses are used for academic and system development purposes only.

Section A: Respondent Information

1. What is your role?
☐ Smallholder Tobacco Farmer ☐ Large-Scale Farmer ☐ Agronomist ☐ Extension Officer ☐ Other: _____
2. How many hectares of tobacco do you grow per season?
☐ Less than 2 ha ☐ 2–5 ha ☐ 5–10 ha ☐ More than 10 ha
3. How many years of farming experience do you have?
☐ Less than 3 years ☐ 3–7 years ☐ 7–15 years ☐ More than 15 years

Section B: Current Monitoring Practices

4. How do you currently monitor your tobacco crop health?
☐ Visual inspection ☐ Agronomist visits ☐ Sensor/equipment ☐ No systematic monitoring
5. How frequently do you inspect your fields?
☐ Daily ☐ Every 2–3 days ☐ Weekly ☐ Fortnightly ☐ Monthly

6. When are stress conditions (water stress, disease) typically detected?
- ☐ Before any visible symptoms ☐ When early symptoms appear ☐ After significant damage is visible ☐ After crop loss has occurred

Section C: Technology Readiness

7. Do you own or have access to a smartphone?
- ☐ Yes, with reliable internet ☐ Yes, with limited internet ☐ No, feature phone only ☐ No mobile phone
8. Are you aware of satellite-based crop monitoring tools?
- ☐ Yes, and I use them ☐ Yes, but I do not use them ☐ I have heard of them ☐ No
9. Would you use a free web-based tool that shows your field's crop health from satellite data?
- ☐ Strongly agree ☐ Agree ☐ Neutral ☐ Disagree ☐ Strongly disagree

Section D: Decision Support Needs

10. What information would be most valuable in a crop monitoring system? (Select all that apply)
- ☐ Water stress level ☐ Disease risk alert ☐ Irrigation schedule ☐ Yield prediction ☐ Field health trend over time ☐ Specific affected zones within the field
11. How important is it to receive farming advice linked to your specific field's satellite data?
- ☐ Very important ☐ Important ☐ Moderately important ☐ Not important
12. What is the biggest challenge in your current monitoring approach?
- ☐ Missing stress until too late ☐ Cost of field inspections ☐ Labour availability ☐ Inaccurate assessments ☐ Other: _____

Thank you for your participation. Your responses directly inform the design of an accessible, farmer-focused satellite monitoring system for Zimbabwean tobacco farming.

Document Analysis: Agritex crop assessment forms, the TIMB annual report (TIMB, 2023), and the Zimbabwe Agricultural Development Strategy II (Ministry of Agriculture, 2022) were reviewed to identify what health parameters are currently recorded and at what frequency, which shaped the analysis output format.

Direct Observation: Attending a field inspection session in Mvurwi district confirmed that a standard visit covers just three to five assessment points in a field — missing an estimated 60–80% of the total area. That observation directly drove the within-field spatial variability requirement.

3.2.2 Technical Feasibility

The technology stack is well-established. The GEE API for Sentinel-2 retrieval is a mature service with thorough documentation, an active developer community, and verified use in Zimbabwean agricultural research (Rwizi et al., 2021). Next.js, React, and scikit-learn are all widely adopted with strong long-term support. Everything in the stack falls within the developer’s existing skill set. No significant technical risks to the project timeline have been identified. Perera et al. (2015) note that systems built on open, well-documented APIs carry substantially lower technical risk than those relying on proprietary or experimental interfaces — a relevant point here.

3.2.3 Hardware Requirements

- **Development machine:** Intel Core i-series or equivalent processor, minimum 8 GB RAM (16 GB recommended), 256 GB SSD storage minimum, running Windows 10/11, macOS, or Ubuntu Linux.
- **Server / cloud:** Vercel or Netlify platform for Next.js deployment (free tier available); managed PostgreSQL service (e.g., Neon, Supabase free tier) for database hosting.
- **End-user device:** Any smartphone or computer with a modern web browser (Chrome, Firefox, Safari) and mobile data or Wi-Fi connectivity.

3.2.4 Software Requirements

- **Operating System:** Windows 11 / macOS 14 / Ubuntu 22.04 LTS
- **Runtime Environments:** Node.js 20 LTS, Python 3.11
- **Web Framework:** Next.js 14 (App Router)
- **Frontend:** React 18, TypeScript 5, React-Leaflet 4, Leaflet.draw
- **Database:** PostgreSQL 16 with PostGIS 3.4 extension
- **ORM:** Prisma 5
- **Machine Learning:** scikit-learn 1.3, pandas 2, matplotlib 3.7

- **Satellite API:** Google Earth Engine Python API (earthengine-api)
- **Weather API:** Open-Meteo (free, no authentication required)
- **Development Environment:** Visual Studio Code
- **Version Control:** Git with GitHub
- **Containerisation:** Docker (optional, for local PostgreSQL development)

3.2.5 Technical Expertise

The development requires working proficiency in TypeScript and React for the frontend; Next.js App Router and REST API design for the backend; Python and scikit-learn for model training; PostgreSQL and PostGIS for spatial data; the GEE API for satellite retrieval; Git for version control; and basic security practices including authentication, password hashing, and HTTPS.

3.3 Economic Feasibility

The economics are favourable from both development and operational perspectives.

3.3.1 Cost-Benefit Analysis

Table 3.1 projects development and operating costs against anticipated savings for a representative cohort of 50 smallholder tobacco farmers over two years.

Table 3.1: Two-Year Cost-Benefit Analysis (USD)

Item	Year 1 (USD)	Year 2 (USD)
Costs		
Development labour (6 months, 1 developer)	2,400	0
Cloud hosting (Vercel + managed PostgreSQL)	0	180
Domain name and SSL certificate	25	25
Testing and contingency	200	100
Total Costs	2,625	305
Benefits (50-farmer cohort)		
Reduced crop losses from early stress detection (est. 5% yield improvement @ USD 2.50/kg average price)	3,125	3,750
Reduced field inspection labour (est. USD 15/visit, 4 visits/season/farmer)	1,500	1,500
Reduced agrochemical waste from targeted treatment (est. USD 20/farmer/season)	500	600
Total Benefits	5,125	5,850
Net Benefit	2,500	5,545
Cumulative Net Benefit	2,500	8,045

The analysis returns a positive net benefit from Year 1, with cumulative net benefit exceeding USD 8,000 over two years across 50 farmers. The 5% yield improvement assumption is deliberately conservative; [Gebbers and Adamchuk \(2010\)](#) show precision agriculture interventions in smallholder contexts typically deliver 5–15% improvements in input efficiency and output.

3.3.2 Tangible Benefits

- **Early stress detection:** A 7–14 day warning before symptoms appear means farmers can irrigate within the effective treatment window, protecting leaf quality and auction prices ([Rouse et al., 1974](#)).
- **Reduced inspection cost:** Automated satellite monitoring reduces the number of physical inspection visits needed, saving farmer time and hired labour.
- **Targeted intervention:** Hotspot identification means agrochemicals can be applied to stressed zones only, rather than the whole field, cutting input costs.
- **Zero subscription cost:** GEE and Open-Meteo are both free. Commercial satellite data for the same coverage would run USD 50–200 per hectare per season.

3.3.3 Intangible Benefits

- **Decision confidence:** Quantitative NDVI data gives farmers an objective basis for timing interventions, reducing the uncertainty that comes with “the crop looks a bit off” judgements.
- **Agricultural literacy:** Regular engagement with trend data builds farmer understanding of plant physiology and precision agriculture over time.
- **Environmental impact:** Targeted irrigation and agrochemical use reduces runoff, groundwater draw-down, and chemical residues.
- **Demonstration effect:** A working open-source deployment in Zimbabwe creates a replicable template for other smallholder crops and other Southern African countries.

3.4 Operational Feasibility

The system fits into existing workflows rather than demanding new ones. Farmers already inspect fields; this just adds a satellite-derived health check they can trigger from a smartphone in under three minutes, requiring nothing beyond the ability to navigate a basic website.

Agritex extension officers are the natural adoption channel. They visit farmers seasonally, understand the agronomic context, and can help with first-time polygon setup. That support structure already exists. [Wieringa \(2014\)](#) argues that operational feasibility hinges not just on whether a system can be used, but whether it’ll be accepted and sustained within existing workflows — and a system designed to complement rather than replace human judgement is far easier to accept.

3.4.1 Schedule Feasibility

The project runs twelve months, March 2024 to February 2025, with clearly defined phases and measurable deliverables at each milestone. Because the entire stack uses open-source frameworks and free APIs, infrastructure setup is fast — which means more of the actual development time goes toward feature work and farmer-facing refinement.

3.5 Work Plan


3.5.1 Work Schedule

The project is organised into five phases:

- 1. **Phase 1 — Research and Requirements (Months 1–2):** Literature review, stakeholder interviews, questionnaire distribution, requirements documentation.
- 2. **Phase 2 — System Design (Month 3):** Architecture design, database schema, API specification, UI wireframing.
- 3. **Phase 3 — Development (Months 4–8):** Database and API implementation, Google Earth Engine integration, machine learning model training, frontend implementation, recommendation engine development.
- 4. **Phase 4 — Testing and Validation (Months 9–10):** Unit testing, integration testing, user acceptance testing with farmer cohort in Mvurwi, model performance evaluation.
- 5. **Phase 5 — Documentation and Submission (Months 11–12):** Project report writing, viva preparation, system deployment documentation.

3.5.2 Gantt Chart

Figure 3.1: Project Gantt Chart (March 2024 – February 2025)



3.6 Conclusion

The requirements baseline is established. Custom build was the only viable path — commercial platforms don’t fit and outsourcing the domain knowledge required isn’t practical.

Requirements for all three user roles are grounded in actual field observations and interviews, not assumptions. Technical, economic, and operational feasibility checks all pass. The Gantt chart sets a realistic twelve-month delivery framework with clear milestones.

Chapter 4: System Analysis

4.1 Introduction

Good system design starts with an honest look at what it's replacing. [Wieringa \(2014\)](#) argues that improvement can't be measured or designed without a clear baseline — and he's right. This chapter starts there: a detailed analysis of the existing tobacco monitoring system, its structural weaknesses, and what the proposed system does differently. The analysis is presented through formal models — context diagrams, data flow diagrams, use case analysis, and UML diagrams.

4.2 Description of the Current System

In Zimbabwe's key tobacco provinces — Mashonaland Central, Mashonaland West, and Manicaland — crop monitoring is manual and entirely inspection-based. The standard cycle works like this:

A farmer or hired inspector walks the field, covering three to five observation points per hectare. At each point, they look at leaf colour, plant posture, and any visible pest or disease signs. Notes go in a notebook or stay in memory. The decision to irrigate, spray, or call an agronomist comes down to experience and what the inspector can see that day.

In areas with Agritex coverage, an extension officer might visit once per season. Cooperative members may receive a seasonal advisory newsletter. Neither is field-specific.

Weather monitoring is informal. Some farmers have a basic rain gauge; temperature is estimated. Zimbabwe's Meteorological Services provides district-level forecasts, but the spatial and temporal resolution is far too coarse to drive individual field management decisions.

None of the smallholder farmers consulted during requirements gathering used any form of satellite, drone, or sensor technology. [Nyoka et al. \(2020\)](#) document the same picture across the sector: remote sensing is almost entirely absent from Zimbabwean smallholder tobacco farming.

4.3 Analysis of the Existing System

The most fundamental limitation is timing. The current system only detects stress once it's visually apparent. [Tucker \(1979\)](#) showed that NDVI begins declining during early water

stress 5–7 days before wilting or leaf colour change becomes visible. That’s the gap: visual inspection is structurally incapable of the early warning that effective intervention requires.


The second problem is coverage. Manual inspection samples a fraction of field area. Localised stress caused by soil compaction, drainage variation, pest entry points, or equipment trampling gets missed. Weiss et al. (2020) demonstrate that satellite-derived within-field NDVI maps routinely expose spatial heterogeneity that aggregate field assessments hide entirely.

The third is continuity. Without systematic record-keeping, there’s no historical baseline to compare current readings against. A declining NDVI trend — arguably the most useful early-warning signal available — is invisible without multi-point historical data. The current system doesn’t collect any.

4.3.1 Context Diagram of the Existing System

Figure 4.1 presents the context diagram of the existing tobacco monitoring system, illustrating the information flows between the farmer and external entities in the current operational environment.

Figure 4.1: Context Diagram — Existing Tobacco Crop Monitoring System



4.3.2 Weaknesses of the Current System

- **Reactive detection:** Stress only registers once symptoms are visible — after the pre-symptom intervention window has already closed.
- **Incomplete spatial coverage:** Manual inspection covers 5–15% of field area per visit. Most of the field goes unmonitored.
- **No data continuity:** No record-keeping means no trend analysis. Slow-developing stress trajectories are invisible.
- **Subjectivity:** Assessments depend on individual inspector experience and vary between observers and visits.

- **Weather isolation:** Environmental conditions aren't integrated with field observations. Stress events aren't correlated with rainfall or temperature anomalies.
- **High opportunity cost:** The labour and time cost of regular manual inspections limits how often monitoring actually happens.

4.4 Description of the Proposed Solution

The proposed system addresses all six weaknesses above through five integrated components:

1. **Interactive Field Definition:** Farmers draw field polygons directly on a satellite basemap using the React-Leaflet interface. Coordinates are stored as GeoJSON in the database.
2. **Live Satellite NDVI Retrieval:** For each analysis request, the system submits the field polygon to GEE, which retrieves all available cloud-free Sentinel-2 Level-2A scenes within a configurable 14-to-30-day window, computes per-pixel NDVI from Band 8 (NIR) and Band 4 (Red), and returns aggregate statistics alongside the spatial NDVI distribution.
3. **Weather Data Integration:** The Open-Meteo API is queried with the field centroid coordinates to retrieve daily temperature and precipitation over the analysis window, which are aggregated into the feature vector.
4. **Machine Learning Classification:** A logistic regression classifier trained on ground-truth satellite and weather observations classifies the field as *Healthy*, *Moderate Stress*, or *High Stress* from five engineered features.
5. **Explainable Recommendations:** A rule-based engine converts the classification and raw feature values into plain-language recommendations, each linked to the specific data observation that triggered it.

4.4.1 Context Diagram and Data Flow Diagrams

Figure 4.2 presents the context diagram for the proposed system, showing the information flows between the system and all external entities.

Figure 4.2: Context Diagram — Proposed Satellite Crop Health Monitoring System





Figure 4.3 presents the Level 1 Data Flow Diagram, decomposing the system into its major processing components.

Figure 4.3: Level 1 Data Flow Diagram — Proposed System



4.5 Requirements Analysis

4.5.1 Functional Requirements

Table 4.1 summarises the thirteen functional requirements derived from user requirements and validated through questionnaire analysis.


Table 4.1: Functional Requirements

FR	Requirement Name	Description
FR01	User Registration	The system shall allow new users to register with name, email, and password, validated against format and uniqueness constraints.
FR02	User Authentication	The system shall authenticate users against stored bcrypt-hashed passwords and issue HTTP-only session cookies on success.
FR03	Field Creation	The system shall enable authenticated farmers to create a field record by drawing a polygon on an interactive map and providing a field name.
FR04	Field Listing	The system shall display all fields belonging to the authenticated user with their names, areas, and last analysis dates.
FR05	Analysis Trigger	The system shall allow farmers to trigger a crop health analysis for any of their fields with a single user action.
FR06	Satellite NDVI Retrieval	The system shall query the Google Earth Engine API with the field polygon and a 14-to-30-day date range to retrieve Sentinel-2 NDVI time-series data.
FR07	Weather Data Retrieval	The system shall query the Open-Meteo API with the field centroid coordinates to retrieve daily temperature and precipitation data for the analysis window.
FR08	Feature Engineering	The system shall compute mean NDVI, NDVI trend slope, NDVI spatial variance, average temperature, and total rainfall from the retrieved satellite and weather data.

FR	Requirement Name	Description
FR09	Health Classification	The system shall apply the trained logistic regression model to the engineered features and classify the field's health status as HEALTHY, MODERATE_STRESS, or HIGH_STRESS.
FR10	Recommendation Generation	The system shall apply four rule-based recommendation rules to the feature values and health classification to generate specific, evidence-linked farming advice.
FR11	Heatmap Visualisation	The system shall render a within-field NDVI heatmap overlay on the field map, highlighting spatial stress hotspots using a colour gradient.
FR12	Analysis History	The system shall store all analysis results and allow users to retrieve a history of analyses for each field.
FR13	Admin User Management	The system shall allow administrators to view, create, update, and deactivate user accounts.

Figure 4.4 presents the Use Case Diagram illustrating actor interactions with the system.

Figure 4.4: Use Case Diagram — Satellite Crop Health Monitoring System



4.5.2 Non-Functional Requirements

- **Performance:** A full analysis — satellite retrieval, feature engineering, ML inference, and recommendations — must complete within 30 seconds under normal net-

work conditions. The web interface should load within 3 seconds on a standard mobile data connection.


- **Security:** Passwords are stored as bcrypt hashes with a minimum work factor of 12. Sessions use HTTP-only cookies to prevent JavaScript-based hijacking. Database row-level constraints ensure users can only access their own field and analysis records.
- **Reliability:** The system must stay available throughout Zimbabwe’s tobacco growing season (October to May). Analysis results are persisted immediately on completion to protect against connectivity interruptions. GEE or Open-Meteo downtime should degrade gracefully with a clear user notification rather than a hard failure.
- **Usability:** A user with secondary school education and basic smartphone literacy should be able to draw a field and trigger an analysis within five minutes, without reading a manual. Health status is conveyed through colour-coded indicators (green, amber, red) that communicate their meaning without accompanying text.
- **Scalability:** The schema supports an unlimited number of fields per user, and PostgreSQL with PostGIS handles spatial queries efficiently as the user base grows. ML inference runs entirely in JavaScript with no Python server, enabling deployment to serverless platforms.
- **Compatibility:** The system must work on the latest two major versions of Chrome, Firefox, Safari, and Edge, on both desktop and mobile, covering Android and iOS.
- **Maintainability:** The ML model is stored as a JSON weight file — update the model by replacing the file, no code change required. Recommendation thresholds are in a separate JSON file for the same reason.

4.6 System Models

4.6.1 UML Activity Diagram

Figure 4.5 presents the UML Activity Diagram for the primary crop health analysis workflow.


Figure 4.5: UML Activity Diagram — Crop Health Analysis Workflow



4.6.2 UML Class Diagram

Figure 4.6 presents the UML Class Diagram for the system's data model as defined in the Prisma schema.


Figure 4.6: UML Class Diagram — System Data Model



4.6.3 UML Sequence Diagram

Figure 4.7 presents the UML Sequence Diagram for the crop health analysis interaction, showing message flows between system components.

Figure 4.7: UML Sequence Diagram — Crop Health Analysis Request



The sequence diagram shows how the full pipeline fits together. The Analysis API does the heavy lifting: it coordinates both external API calls, runs the feature engineering and ML inference internally, writes the result to the database, and returns the formatted response to the web client in a single request cycle.

4.7 Conclusion

Six weaknesses in the existing system were identified and all six have a direct answer in the proposed design. The formal models produced in this chapter — context diagrams, Level 1

DFD, use case diagram, activity diagram, class diagram, and sequence diagram — collectively define the architecture that guides implementation. Four capabilities distinguish the proposed system from what it replaces: interactive field boundary definition, live satellite NDVI retrieval with temporal trend analysis, machine learning health classification, and explainable rule-based recommendations. Together they turn a reactive, inspection-based monitoring regime into a proactive, data-driven one.

Chapter 5: System Design

5.1 Introduction

The analysis work in Chapter 4 established what the system must do — the actors, the data flows, the use cases, and the requirements that bound the design. This chapter records the decisions about *how* those requirements will be satisfied: what layered structure the application follows, how the database is organised to support both operational and longitudinal use, what the machine learning pipeline looks like from raw satellite observation to health classification, and what each page of the interface presents to a farmer. Getting the architecture right at this point matters because, as [Sommerville \(2016\)](#) notes, design errors are the most expensive to fix later, since every implementation phase builds directly on top of them.

The overarching design philosophy is modular separation. Each layer has one job, exposes a clean interface to the layers around it, and does not reach into the internals of anything else. That principle runs through every decision in this chapter — from the way GEE and Open-Meteo calls are isolated in their own modules, to the way Python model training is entirely separate from the JavaScript inference that runs at deployment time. Keeping responsibilities narrow is what makes a system like this maintainable by a single developer and extensible in the future without dismantling what already works.

5.2 System Design

Overall Design Approach

The system follows a four-layer modular architecture. Satellite data retrieval, backend coordination, intelligence processing, and browser-based presentation each occupy their own layer. Communication between layers flows through defined interfaces rather than direct internal access.

- **Satellite and weather data acquisition layer**
 - This layer's only job is to accept a GeoJSON polygon and a date range and return a structured feature set. It holds no business logic and nothing about users or fields.
 - Google Earth Engine handles the satellite side. The COPENICUS/S2_SR_HARMONIZED collection is queried with the field polygon as the region of interest, filtered to

scenes with cloud cover below 20% using the QA60 band, and sampled over the analysis window. Per-pixel NDVI is computed at 10-metre resolution from Band 8 (NIR) and Band 4 (Red): $NDVI = (B_8 - B_4) / (B_8 + B_4)$. The spatial pixel array is retained for heatmap rendering. ([ESA, 2023](#); [Gorelick et al., 2017](#))

- Weather data is retrieved in parallel from the Open-Meteo API using the field centroid coordinates. Daily temperature and precipitation records over the same window are aggregated into average temperature and total rainfall. Running these two API calls concurrently reduces overall analysis latency. ([Open-Meteo, 2023](#))
- The layer returns five engineered values plus a pixel-level NDVI grid. Nothing that calls this layer needs to know anything about HTTP requests, GEE authentication, or how NDVI is computed.

- **Backend processing and storage layer**

- Next.js API routes receive incoming requests, validate the session and ownership, orchestrate calls to the acquisition and intelligence layers, and persist results to the database. Validation is the first thing that runs on every request — an unauthenticated or unauthorised request never reaches any computation.
- Storing raw feature values alongside derived outputs (health status, risk flags) is a deliberate design choice. It means a stored analysis record can be re-evaluated if the model is retrained, and the reasoning behind any past recommendation can be traced back to the data that produced it. ([Prisma, 2023](#))
- The route handler itself is deliberately thin. The satellite module, inference module, and recommendation module each live in separate files, so any one of them can be replaced or updated without touching the request-handling logic.

- **Intelligence layer**

- The logistic regression classifier runs entirely in JavaScript using the exported JSON weight file. No Python interpreter is needed at runtime — the model coefficients and intercepts are loaded once at server start and reused across all inference requests. ([Pedregosa et al., 2011](#))
- Classification returns one of HEALTHY, MODERATE_STRESS, or HIGH_STRESS. The NDVI trend slope is separately classified as IMPROVING, STABLE, or DECLINING from a threshold applied to the raw slope value.
- The recommendation engine evaluates four independent rules against the feature values and health status. Rules are not mutually exclusive: a field with low NDVI, low rainfall, high variance, and elevated temperature fires all four applicable rules simultaneously. ([Liakos et al., 2018](#))

- **Presentation layer**

- The browser application handles all farmer-facing interaction through server-rendered Next.js 14 pages with TypeScript. Registration, login, field creation, analysis triggering, and result review are each their own page, making the interface navigable even on slow mobile connections.
- The interactive map uses React-Leaflet with an ESRI World Imagery satellite tile layer as the basemap, so farmers can draw polygon boundaries directly over recognisable aerial imagery of their own land. The Leaflet-Draw plugin provides the drawing tools, restricted to the polygon type only.
- Health status is communicated primarily through colour — green for Healthy, amber for Moderate Stress, red for High Stress. [Rose et al. \(2016\)](#) identifies plain, immediately readable outputs as one of the strongest predictors of whether farmers actually use a decision support tool. That finding shaped the entire results page design.

5.2.1 How will the System Work?

The main analysis workflow runs from the farmer selecting a field through to a complete results page being displayed. The steps below trace that flow in the order it actually executes.

- **Authentication and session management**

- The farmer logs in with email and password. The backend validates the submitted password against the stored bcrypt hash and, on success, issues an HTTP-only session cookie. Every API request thereafter is checked against that cookie before any processing begins.

- **Field creation with interactive boundary drawing**

- On the Create Field page, a satellite map loads centred on Zimbabwe. The farmer enters a field name and activates the polygon drawing tool. Clicking around the field perimeter on the satellite imagery closes the polygon; it can be edited or redrawn before saving.
- On submission, the GeoJSON polygon is sent to `/api/field`. The backend computes the field area using the Shoelace formula, reverse-geocodes the centroid to a place name through the OpenStreetMap Nominatim API, and stores the record.


- **Satellite data retrieval**

- When the farmer triggers an analysis, the stored polygon is retrieved and passed to the satellite module. The GEE service account authenticates and queries CLOUDMANAGER/S2_SR_HARMONIZED over a 30-day lookback window, returning the per-pixel NDVI series and spatial grid. The Open-Meteo call runs in parallel.
- **Feature engineering, classification, and recommendation**
 - Mean NDVI, trend slope, within-field variance, average temperature, and total rainfall are assembled into the feature vector. The logistic regression classifier scores each class and returns the highest-scoring label. The recommendation engine then evaluates all four rules and returns the matched advice strings.
- **Storage and display**
 - The complete analysis record is written to the database. The formatted response — health status, feature values, risk flags, recommendations, and heatmap pixel array — is returned to the browser. The results page renders the health badge, recommendation cards, environment summary, NDVI heatmap overlay, and the NDVI history chart.

5.3 Solution Architecture

Figure 5.1 presents the proposed system architecture. The four internal layers stack vertically, with the two external data services — Google Earth Engine and Open-Meteo — connecting into the application layer from the right. All farmer interaction flows through the top presentation layer, and all persistence flows down to the data layer.

Figure 5.1: Proposed System Architecture — TobaccoGuard



The analysis route in the application layer is the most complex component in the architecture. It is the only point where both external API calls are made, it is the only caller of

the ML inference and recommendation functions, and it is responsible for persisting the complete result before returning the response to the browser. Everything else in the system is kept simple by design.

5.4 Database Modelling

Purpose of the Database

The database serves two roles. First, it is the operational backbone of the application — it stores user credentials, field polygons, and analysis results in a way that persists across sessions and across seasons. Second, it is the foundation for longitudinal monitoring: every analysis run for a field is stored as a distinct record with a timestamp, which means the system can surface historical NDVI trend data alongside any current reading. That time-series view is, in many ways, more useful to a farmer than any single analysis in isolation. ([Wolfert et al., 2017](#))

- **Choice of database technology**

- SQLite was selected over a server-managed database for the prototype. Its file-based architecture removes the deployment overhead of a separate database process, and Prisma’s ORM layer abstracts the SQL entirely, making a future migration to PostgreSQL a single provider setting change. For a bounded prototype user base without high concurrent write volume, SQLite is sufficient. ([Prisma, 2023](#))


- **Core datasets**

- The most critical data the system stores is the GeoJSON polygon for each field — it is the input to every analysis and cannot be reconstructed from any other record. Analysis records are the second most important: they store both the raw feature values and the derived outputs, so results are auditable and the model’s reasoning can be traced to specific observations.
- The three-tier hierarchy — User owns Fields, Fields accumulate Analyses — is the central relational structure. Cascade deletes flow downward: removing a field removes all its historical analyses; removing a user removes all their fields and analyses.

5.4.1 E-R Diagram

Figure [5.2](#) shows the entity-relationship diagram for the database. The three entities correspond directly to the three Prisma models in the schema.

Figure 5.2: Entity-Relationship Diagram — TobaccoGuard Database



5.4.2 Data Dictionary

The following tables define the fields, data types, and constraints for each entity in the database.

USER (Farmer account records)

Field	Data Type	Description	Constraints
id	TEXT	Unique identifier for the user	PK, unique, not null
name	TEXT	Farmer's full name as entered during registration	unique, not null
email	TEXT	Farmer's email address, used for login	unique, not null
password	TEXT	bcrypt hash of the user's password, work factor 12	not null
createdAt	DATETIME	Timestamp of account creation, set automatically	not null

FIELD (Tobacco field records)

Field	Data Type	Description	Constraints
id	TEXT	Unique identifier for the field record	PK, unique, not null
name	TEXT	Field name as entered by the farmer	not null
cropType	TEXT	Crop grown in the field; defaults to Tobacco	default: Tobacco
polygon	TEXT	GeoJSON representation of the drawn field boundary	not null
area	REAL	Calculated field area in hectares using Shoelace formula	nullable
location	TEXT	Reverse-geocoded place name from field centroid	nullable
userId	TEXT	Reference to the owning user's account	FK → USER.id
createdAt	DATETIME	Timestamp of field record creation	not null

ANALYSIS (Crop health analysis records)

Field	Data Type	Description	Constraints
id	TEXT	Unique identifier for the analysis run	PK, unique, not null
fieldId	TEXT	Reference to the analysed field	FK → FIELD.id
meanNDVI	REAL	Mean NDVI across all cloud-free pixels in the analysis window	not null
ndviTrend	ENUM	Temporal NDVI trajectory: IMPROVING, STABLE, or DECLINING	not null
healthStatus	ENUM	ML classifier output: HEALTHY, MODERATE_STRESS, or HIGH_STRESS	not null
avgTemperature	REAL	Average daily temperature (°C) over the analysis window	not null

Field	Data Type	Description	Constraints
totalRainfall	REAL	Cumulative precipitation (mm) over the analysis window	not null
waterStressRisk	BOOL	True if rainfall < 10 mm AND mean NDVI < 0.45	not null
diseaseRisk	BOOL	True if NDVI spatial variance > 0.05	not null
rawData	TEXT	JSON string storing raw feature values and data source metadata	not null
recommendations	TEXT	JSON array of plain-language recommendation strings	not null
heatmapData	TEXT	JSON array of [lat, lng, ndvi] triples for heatmap rendering	nullable
createdAt	DATETIME	Timestamp of analysis run	not null

5.4.3 Database Schema

The Prisma schema below defines all three models, the two enumerations, and the relationship constraints. The complete schema file is reproduced in full in [Appendix E](#).

Listing 5.1: Prisma schema – core models (abridged)

```

model User {
  id          String    @id @default(cuid())
  name        String
  email       String    @unique
  password    String
  fields      Field[]
  createdAt   DateTime  @default(now())
}

model Field {
  id          String    @id @default(cuid())
  name        String
  cropType    String    @default("Tobacco")
  polygon     String    // GeoJSON stored as JSON string
  area        Float?
  location    String?
  userId      String

```

```

user      User      @relation(fields: [userId], references: [
    id])
analyses  Analysis[]
createdAt DateTime  @default(now())
}

model Analysis {
    id          String      @id @default(cuid())
    fieldId     String
    field       Field       @relation(fields: [fieldId],
        references: [id])
    meanNDVI    Float
    ndviTrend   NDVITrend
    healthStatus HealthStatus
    avgTemperature Float
    totalRainfall Float
    waterStressRisk Boolean
    diseaseRisk Boolean
    rawData     String
    recommendations String
    heatmapData String?
    createdAt   DateTime    @default(now())
}

enum NDVITrend { IMPROVING STABLE DECLINING }
enum HealthStatus { HEALTHY MODERATE_STRESS HIGH_STRESS }

```

5.5 Algorithm Design

The analysis algorithm converts a GeoJSON polygon into a health classification and a set of farming recommendations through a six-step pipeline: ingest and validate → satellite and weather retrieval → feature engineering → ML inference → post-processing and risk flagging → storage and response. Each step is described below. (Liakos et al., 2018)

5.5.1 Inputs to the Algorithm

- The algorithm receives a field record containing:
 - The GeoJSON polygon geometry (a FeatureCollection wrapping a single Polygon feature).

- The `fieldId` (used to verify ownership and store results).
- The authenticated user's session identifier.
- Derived from the polygon during the analysis run:
 - Sentinel-2 Level-2A per-pixel NDVI values over a 30-day window, drawn from the COPENNICUS/S2_SR_HARMONIZED GEE collection.
 - Daily temperature and precipitation for the same window from Open-Meteo.

5.5.2 Pre-processing and Validation

- **Request validation**
 - The session cookie is checked. If absent or invalid, the request is rejected with HTTP 401 before any computation begins.
 - The field record is fetched from the database filtering on both `fieldId` and `userId`. A record that doesn't belong to the authenticated user returns HTTP 404 rather than HTTP 403, to avoid confirming the existence of other users' fields.
- **Polygon parsing**
 - The stored polygon JSON string is parsed. If parsing fails, the request is rejected. The polygon coordinates are extracted and validated for minimum vertex count (at least 4, including the closing coordinate) before being submitted to GEE.
- **Satellite data quality**
 - GEE scenes are filtered to cloud cover below 20%. If fewer than two qualifying scenes exist in the 30-day window, the lookback is extended to 60 days. If still insufficient, the analysis proceeds with a data quality flag in the raw data JSON.

5.5.3 Feature Engineering

The five features used by the logistic regression model are engineered from the raw satellite and weather observations as follows.

- **Mean NDVI** (`mean_ndvi`): The arithmetic mean of all per-pixel NDVI values across all qualifying scenes in the analysis window. Pixels with NDVI below -0.1 (water, shadow) are excluded from the mean computation.

- **NDVI trend slope** (`ndvi_trend`): A linear regression slope fitted to the sequence of per-scene mean NDVI values ordered by acquisition date. A positive slope indicates improving vegetation; a negative slope indicates decline. The slope is classified as IMPROVING if > 0 , DECLINING if < -0.01 , and STABLE otherwise. (Weiss et al., 2020)
- **NDVI variance** (`ndvi_variance`): The spatial variance of per-pixel NDVI values within the field polygon, computed from the pixel distribution in the most recent cloud-free scene. High variance indicates irregular within-field vegetation patterns, which may indicate patchy disease, pest damage, or waterlogging in isolated zones. (Nyoka et al., 2020)
- **Average temperature** (`avg_temperature_c`): Mean of daily maximum temperatures over the analysis window from Open-Meteo.
- **Total rainfall** (`total_rainfall_mm`): Sum of daily precipitation over the same window.

5.5.4 Model Selection and Responsibilities

A multinomial logistic regression classifier is used for the three-class crop health problem. The choice was informed by Liakos et al. (2018), who reviewed forty machine learning studies in agriculture and found that logistic regression frequently matches more complex models when features are well-engineered and class boundaries are reasonably stable. The key advantages here are interpretability — the coefficient magnitudes show which features most influence each class — and deterministic inference, since there is no randomness in the forward pass at prediction time.

- **Responsibility 1 – Health classification:** Assigns the field to HEALTHY, MODERATE_STRESS, or HIGH_STRESS from the five-feature vector.
- **Responsibility 2 – Risk flagging:** Two binary flags are derived separately from the raw features using threshold rules. `waterStressRisk` fires when `totalRainfall` < 10 mm and `meanNDVI` < 0.45 . `diseaseRisk` fires when `ndviVariance` > 0.05 .
- **Responsibility 3 – Trend tracking:** The NDVI trend direction is stored alongside the health status as a secondary indicator for longitudinal review.

5.5.5 Inference and Post-processing Logic

1. **Score computation:** For each class j , a linear score is computed: $z_j = \mathbf{w}_j \cdot \mathbf{x} + b_j$, where $\mathbf{w}_j$ is the row of coefficients for class j , $\mathbf{x}$ is the feature vector, and b_j is the class intercept.

2. **Class selection:** The class with the highest score is selected as the prediction: $\hat{y} = \arg \max_j(z_j)$.
3. **Recommendation generation:** The four recommendation rules are evaluated independently:
 - **R1 – Water stress:** $\text{NDVI} < 0.45$ and $\text{rainfall} < 10 \text{ mm} \rightarrow$ irrigation recommended.
 - **R2 – Declining moderate:** $\text{trend} = \text{DECLINING}$ and $0.45 \leq \text{NDVI} \leq 0.60 \rightarrow$ monitor closely.
 - **R3 – Stable healthy:** $\text{NDVI} > 0.60$ and $\text{trend} = \text{STABLE}$ or $\text{IMPROVING} \rightarrow$ no action required.
 - **R4 – Disease signal:** $\text{variance} > 0.05$ and $\text{temperature} > 25^\circ\text{C} \rightarrow$ inspect for disease or pests.
4. **Heatmap generation:** The retained per-pixel NDVI grid is converted to a list of [lat, lng, intensity] triples, where intensity = $1 - \text{normalised NDVI}$ (so that red corresponds to low NDVI / high stress, and green to high NDVI / healthy). Pixels are sampled at fixed intervals to keep the payload manageable.

5.5.6 Algorithm Pseudocode

```

Receive analysis request with fieldId and session cookie
Validate session cookie and ownership of fieldId
Retrieve field polygon from database
Parse GeoJSON polygon and extract coordinate array
Submit polygon to Google Earth Engine
  Filter COPERNICUS/S2_SR_HARMONIZED by polygon and 30-day window
  Apply QA60 cloud mask (< 20% cloud cover)
  Compute per-pixel NDVI = (B8 - B4) / (B8 + B4)
  Return NDVI time-series and spatial pixel grid
In parallel, request weather data from Open-Meteo
  For field centroid coordinates over same 30-day window
  Return daily temperature and precipitation
Compute engineered features:
mean_ndvi      = mean of all qualifying pixels
ndvi_trend     = linear regression slope over time-series
ndvi_variance  = spatial variance of pixel distribution
avg_temperature = mean of daily maximum temperatures
total_rainfall = sum of daily precipitation

```



```

Load JSON model weights (coefficients + intercepts)
For each class: compute  $z = \text{dot}(\text{weights}[j], \text{features}) + \text{bias}[j]$ 
Select class with highest  $z$  as health status
Evaluate four recommendation rules against features and status
Generate heatmap point array from pixel grid
Write complete analysis record to database
Return response to browser (status, features, recommendations, heatmap)

```

5.6 Interface Design

The interface was designed around one principle: the primary output of any page should be readable in under three seconds, without scrolling, by a farmer with basic digital literacy. That constraint pushed towards large, colour-coded status indicators, minimal text on the primary views, and details moved into expandable sections for users who want them. [Rose et al. \(2016\)](#)

5.6.1 Login and Registration Interface

The Login and Sign-Up pages are deliberately simple. The farmer's task here is purely credential entry — the design should not distract from that with navigation or secondary content.

Figure 5.3: Wireframe — Login Page

TobaccoGuard

Sign in to your account
 Monitor your tobacco fields with satellite data

Email address

Password

Sign In


Don't have an account? [Create one here](#)

Free to use • No subscription fees • Powered by Google Earth Engine

5.6.2 Dashboard Interface

The Dashboard is the farmer's landing page after login. Its purpose is to surface the health status of every field at a glance, with a clear path to trigger a new analysis or view prior results.

Figure 5.4: Wireframe — Dashboard (Field List View)



5.6.3 Create Field Interface

The Create Field page uses a split layout: a form panel on the left and an interactive satellite map on the right. The map occupies most of the viewport so that drawing a polygon over recognisable aerial imagery is as natural as possible.

Figure 5.5: Wireframe — Create Field (Interactive Map View)

TobaccoGuard | Create New Field

Field Details

Field Name

Drawing Instructions

1. Click the polygon tool (top right of map)
2. Click around your field boundary
3. Click the first point to close the shape
4. Edit or redraw as needed
5. Click Save Field below

Navigate to location

Go

Computed Area

– ha

Save Field

Cancel

Lat: -17.8300
Lng: 31.0500


Draw Polygon

Satellite Imagery)

5.6.4 Analysis Results Interface

The Analysis Results page is the primary output surface of the system. It is designed to present the most important information — the health status and the recommended actions — without requiring the farmer to scroll or navigate further.

Figure 5.6: Wireframe — Analysis Results Page



5.7 Conclusion

The design stage produced a complete specification of the system before a line of implementation code was committed. The four-layer modular architecture separates satellite data retrieval, backend coordination, machine learning intelligence, and browser presentation into components that can be built, tested, and updated independently. The database design supports both the operational and longitudinal requirements of the system through three clearly related entities. The algorithm design provides a step-by-step pathway from field polygon to health classification and spatial heatmap. And the wireframes translate the functional requirements into concrete page layouts that prioritise the farmer’s primary information needs. Chapter 6 takes these design decisions and records how each one was realised in working code.

Chapter 6: System Implementation

6.1 Introduction

This chapter records how the designs from Chapter 5 were turned into a working application. It covers the development environment, the application’s directory structure, and the key implementation decisions for each major component: the interactive map and frontend, the backend API routes, the Google Earth Engine satellite integration, the machine learning pipeline, and the database. Where a full code listing is given in the appendices, this chapter focuses on the design rationale rather than reproducing every line. The intent is to explain *why* things were built the way they were, not just *what* was built.

The implementation was carried out over roughly four months, working iteratively rather than phase-by-phase. The machine learning model was trained and exported early so that the full analysis pipeline could be exercised end-to-end from the start, catching integration issues early rather than at the end. Interface work and backend work progressed in parallel, connected through an agreed API contract that was finalised during the design phase.

6.2 Development Environment and Tools

Table 6.1 lists the core tools used during development, along with the version and the reason each was chosen.

Table 6.1: Development Tools and Technologies

Tool / Technology	Version	Rationale
Next.js	14	Full-stack React framework with API routes, server rendering, and serverless deployment support. Single language (TypeScript) across frontend and back-end.
TypeScript	5.x	Static typing catches integration errors between modules at compile time, which matters when passing feature vectors between the satellite module, inference module, and API route.
Tailwind CSS	4.x	Utility-first CSS framework enabling rapid UI iteration without writing a separate stylesheet.
shadcn/ui	Latest	Accessible, unstyled Radix-based component library — badges, cards, accordions, and tooltips used throughout the results page.
React-Leaflet	4.x	React bindings for Leaflet.js, providing the interactive satellite map component.
Leaflet-Draw	1.x	Leaflet plugin providing the polygon drawing toolbar and shape event callbacks.
Leaflet.heat	0.2	Lightweight heatmap canvas overlay for Leaflet, used to render the per-pixel NDVI stress layer.
Recharts	2.x	React charting library used for the NDVI history area chart on the results page.
Prisma ORM	5.x	Schema-first ORM with TypeScript type generation, SQLite support, and straightforward migration tooling. (Prisma, 2023)
SQLite	3.x	Embedded, file-based relational database for the prototype. Zero external service dependency.
Python	3.11	Used exclusively for the offline model training pipeline; not required at application runtime.
scikit-learn	1.3	Provides the LogisticRegression classifier used for training. (Pedregosa et al., 2011)
bcryptjs	2.x	Password hashing library with configurable work factor; used for user authentication.

6.3 Application Structure

The application follows the Next.js App Router convention. Table 6.2 gives an overview of the directory layout and the purpose of each major folder.

Table 6.2: Application Directory Structure

Directory / File	Purpose
app/	Next.js App Router root. Contains all pages and API routes.
app/api/auth/	Login, sign-up, and logout API route handlers.
app/api/field/	Field creation and deletion API route handlers.
app/api/analyse/	Core analysis pipeline route handler.
app/dashboard/	Dashboard page showing the user's field list.
app/field/create/	Field creation page with interactive map.
app/field/[fieldId]/analyse/	Analysis results page.
components/Map/	Leaflet map components (SSR wrapper, draw map, view map).
components/ui/	shadcn/ui component library (badges, cards, buttons, etc.).
lib/gee/	Google Earth Engine integration module.
lib/ml/	JavaScript logistic regression inference module.
lib/analysis/	Recommendation engine.
lib/geo/	Polygon utility functions (area, centroid, GeoJSON parsing).
models/src/	Python model training and export scripts.
models/models/	Exported JSON model weights and threshold configuration.
prisma/	Prisma schema file and SQLite database file.

6.4 Frontend Implementation

6.4.1 Interactive Map Component

The interactive satellite map is implemented in `components/Map/LeafletMap.tsx`. Because Leaflet requires access to the browser's window object, it cannot run during Next.js server-side rendering. The component is loaded through `next/dynamic` with `ssr: false`, wrapped in a lightweight `FieldMap.tsx` component that renders a loading placeholder during the hydration phase.

The map configuration enables two tile layers simultaneously: ESRI World Imagery for the satellite basemap and a Stamen Toner Labels layer at 70% opacity for place name overlay. Farmers can see both the satellite imagery and recognisable place names when

navigating to their field area.

Polygon drawing is handled by the `EditControl` component from the `react-leaflet-draw` library. All draw tools except the polygon tool are disabled in the control configuration. When a polygon is created, edited, or deleted, the corresponding event handler serialises the shape to GeoJSON and passes it up through a callback prop. The polygon is styled with a 40% opacity emerald fill so that the satellite basemap remains visible through the drawn boundary.

6.4.2 Analysis Results Page

The results page (`app/field/[fieldId]/analyse/page.tsx`) triggers the analysis on mount using a `useEffect` hook, posting the field identifier to `/api/analyse` and rendering a loading spinner while the analysis runs. On response, the health status badge, recommendation cards, risk indicators, environment summary, heatmap, and NDVI history chart are all rendered from the returned JSON.

The NDVI heatmap is rendered as a `Leaflet.heat` layer added imperatively to the map instance using `useMap()` inside a dedicated `HeatmapLayer` component. The gradient is configured to map low intensity (high NDVI, low stress) to green and high intensity (low NDVI, high stress) to red, so the spatial colour coding aligns with the health badge colours used elsewhere in the interface.

The NDVI history area chart uses `Recharts AreaChart` with a linear gradient fill. The five data points represent the per-scene mean NDVI values across the analysis window, plotted in order of acquisition date.

6.5 Backend API Implementation

6.5.1 Authentication Routes

Session management is handled through HTTP-only cookies rather than `localStorage`-based JWT tokens. An HTTP-only cookie cannot be read by client-side JavaScript, which prevents a whole category of token-theft attacks. The cookie stores the user's `id` from the database; every protected API route reads that cookie and uses it to scope database queries to the authenticated user's records only.

Password hashing uses `bcryptjs` with a work factor of 12. At work factor 12, a single hash computation takes approximately 250–300 ms on typical server hardware — slow enough to make brute-force attacks impractical, but fast enough that login latency remains acceptable.

6.5.2 Field Management Routes

The `/api/field` POST handler validates the incoming field name, parses the GeoJSON polygon, computes the field area in hectares using the Shoelace formula, and calls the OpenStreetMap Nominatim API to reverse-geocode the polygon centroid to a human-readable location name. All three computed values — area, location, and GeoJSON — are stored in the field record.

The Shoelace formula computes the area of a polygon from its vertex coordinates:

$$A = \frac{1}{2} \left| \sum_{i=0}^{n-1} (x_i y_{i+1} - x_{i+1} y_i) \right| \quad (6.1)$$

The coordinates are given in decimal degrees, so the result is scaled from square degrees to square kilometres using the Earth’s mean radius (6371 km) and then converted to hectares.

6.5.3 Analysis Route

The `/api/analyse` POST handler is the most involved route in the application. It performs the following steps in order:

1. Validates the session cookie and fetches the field record, filtering on both `fieldId` and `userId` to prevent cross-user access.
2. Parses the stored GeoJSON polygon and extracts the coordinate array.
3. Calls `getSatelliteData(polygon)` from the GEE module. Inside that function, the GEE call and the Open-Meteo call are dispatched with `Promise.all()` so they run concurrently.
4. Assembles the five-feature vector from the returned satellite and weather data.
5. Calls `predictCropHealth(features)` from the inference module.
6. Calls `getRecommendations({...})` from the recommendation engine.
7. Writes the complete analysis record to the database using `Prisma’s analysis.create()`.
8. Returns the formatted response including all feature values, health status, risk flags, recommendations, and heatmap data.

6.6 Google Earth Engine Integration

6.6.1 Authentication and Service Account

The GEE integration authenticates using a Google service account. The service account private key and project ID are stored as environment variables (`GEE_PRIVATE_KEY`, `GEE_SERVICE_ACCOUNT`, `GEE_PROJECT`), loaded at server start, and used to initialise the GEE client. The service account is registered under a non-commercial research project, qualifying for the free 150 EECU-hours per month Community tier. (Gorelick et al., 2017)

6.6.2 NDVI Retrieval Pipeline

The satellite data module implements the following retrieval pipeline for each analysis request:

1. The field's GeoJSON polygon is converted to a GEE Geometry object.
2. The COPERNICUS/S2_SR_HARMONIZED image collection is filtered to the field geometry and a 30-day lookback window. Scenes with more than 20% cloud cover are excluded using the `CLOUDY_PIXEL_PERCENTAGE` property.
3. For each qualifying scene, per-pixel NDVI is computed as a new image band: `ndvi = image.normalizedDifference(['B8', 'B4'])`.
4. The NDVI images are reduced to a temporal mean using `ImageCollection.mean()`, and separately to a time-series of per-scene mean values for trend computation.
5. The spatial pixel distribution within the polygon is sampled at 10-metre resolution using `Image.sample()` to produce the within-field variance and the heatmap pixel array.

6.6.3 Heatmap Data Generation

The per-pixel NDVI sample from GEE returns a `FeatureCollection` where each feature represents one 10-metre pixel and carries its geographic coordinates and NDVI value. This collection is converted on the server into a compact array of `[latitude, longitude, intensity]` triples, where intensity is $(1 - \text{NDVI})$ clipped to $[0, 1]$ — so that a pixel with $\text{NDVI} = 0.9$ (healthy) has intensity 0.1 (green on the heatmap) and a pixel with $\text{NDVI} = 0.2$ (stressed) has intensity 0.8 (red on the heatmap).

The pixel array is downsampled at a fixed interval before serialisation to keep the heatmap JSON payload under 50 KB for typical field sizes. The full pixel resolution is retained in the raw data JSON for potential future use.

6.7 Machine Learning Implementation

6.7.1 Training Pipeline

The logistic regression classifier is trained entirely offline in Python using scikit-learn 1.3. The training script (`models/src/train_health_model.py`) is reproduced in full in Appendix A. A summary of the process:

1. A labelled dataset of 610 observations is loaded from `models/data/raw/dataset.csv`. Each row contains five feature values and a health status label (HEALTHY, MODERATE_STRESS, or HIGH_STRESS), assigned using the agronomically validated NDVI thresholds from Table 7.2.
2. An 80/20 stratified train-test split is applied with `random_state=42` for reproducibility, giving 488 training samples and 122 test samples.
3. A `LogisticRegression` model is fitted with `max_iter=1000` to ensure convergence.
4. The trained model is evaluated on the held-out test set, achieving 97% overall accuracy.

6.7.2 Model Export and JavaScript Inference

After training, the model weights are exported to a JSON file using the export script (`models/src/export_model.py`, Appendix B). The JSON file stores the coefficient matrix and intercept vector:

Listing 6.1: Structure of the exported model JSON

```
{
  "classes":      ["HEALTHY", "HIGH_STRESS", "MODERATE_STRESS"],
  "coefficients": [[...], [...], [...]], // shape: [3, 5]
  "intercept":   [...]                  // length: 3
}
```

The JavaScript inference module (`lib/ml/inference.ts`, Appendix C) imports this JSON file directly. At runtime it computes a linear score for each class and returns the class with the highest score. The complete inference function is seven lines of TypeScript. No external ML library is needed at runtime — the exported weights are just numbers.

6.7.3 Recommendation Engine

The recommendation engine (`lib/analysis/recommendations.ts`, Appendix D) takes the health status, NDVI trend direction, and raw feature values as inputs and returns an array of recommendation strings. The four rules from Table 7.4 are evaluated in sequence. Rules are independent and additive — all rules that fire are included in the output, not just the first one. A default recommendation is returned if no rule fires, ensuring the results page always has at least one item to display.

6.8 Database Implementation

The Prisma schema (Appendix E) defines all three models and the two enumerations. The development database is a single SQLite file (`dev.db`) stored at the project root. The Prisma client is instantiated once in `lib/prisma.ts` using a global singleton pattern to avoid opening multiple database connections during Next.js hot-reload in development.

Database migrations are managed through Prisma Migrate. Running `npx prisma migrate dev` reads the schema, computes the diff against the current database state, generates a SQL migration file, and applies it — all in a single command. For the production transition to PostgreSQL, the only required change is the `datasource db { provider = "postgresql" }` line in `schema.prisma`. (Prisma, 2023)

6.9 Conclusion

All seven technical objectives from Chapter 1 have been addressed in the implementation. The full-stack Next.js application is deployable serverlessly, the logistic regression classifier runs in JavaScript without a Python backend, GEE retrieves real Sentinel-2 NDVI data for the user's drawn polygon, and the within-field heatmap renders per-pixel stress levels directly on a Leaflet map. Chapter 7 evaluates the implemented system quantitatively — measuring the classifier's performance on held-out data, assessing each functional requirement against the implementation, and testing the non-functional properties including performance and security.

Chapter 7: Testing, Evaluation and Conclusion

7.1 Introduction

This chapter closes the project report. It reviews what was actually built against the objectives set in Chapter 1, evaluates the machine learning model against quantitative benchmarks, tests the system against the functional and non-functional requirements defined in Chapter 3, acknowledges the prototype’s real limitations honestly, and proposes the most valuable directions for future work. It also puts the project’s contribution in context — within the precision agriculture landscape and within Zimbabwe’s agricultural development ambitions.

7.2 Summary of Project Achievements

The project produced a complete, deployable web application with six integrated components:

1. **Interactive Web Application:** A full-stack Next.js 14 application with TypeScript, covering farmer registration and authentication, interactive field polygon drawing via React-Leaflet, on-demand health analysis, result visualisation, and analysis history. It runs in any modern browser on desktop or mobile without installation.
2. **Satellite Data Integration Layer:** A geospatial retrieval module that acquires per-field NDVI estimates from satellite imagery, computes mean NDVI, temporal trend slope, and within-field variance, and returns a structured feature set for downstream classification.
3. **Weather Data Integration:** A live connection to the Open-Meteo API that pulls 14-day historical daily temperature and precipitation records for each field’s centroid, giving the model the environmental context it needs for a robust assessment.
4. **Logistic Regression Classifier:** A crop health model trained in Python with scikit-learn on 610 labelled observations, exported as a JSON weight file, and deployed for inference entirely in JavaScript inside the Next.js app. No Python server is required at runtime.
5. **Explainable Recommendation Engine:** A four-rule decision system that converts the classifier output and raw feature values into plain-language advice on irrigation,

disease monitoring, stable crop management, and pest inspection — each recommendation tied to the specific data reading that triggered it.

6. **Persistent Database Layer:** A Prisma ORM database backend storing all user, field, and analysis records, with full per-field history to support longitudinal trend review across seasons.

7.3 Evaluation of Technical Objectives

Table 7.1 evaluates the degree to which each of the seven technical objectives stated in Chapter 1 has been achieved in the implemented system.

Table 7.1: Evaluation of Technical Objectives Against Implementation

Obj.	Objective Statement	Implementation Outcome	Status
O1	Design and implement a web-based platform enabling farmers to define field boundaries and trigger analysis on demand.	Next.js 14 application with React-Leaflet polygon drawing, single-action analysis trigger, and full result display implemented and deployable.	Achieved
O2	Integrate live Sentinel-2 Level-2A imagery through GEE to compute real NDVI for user-defined polygons.	A satellite data module computing NDVI, trend, and variance from field coordinates is implemented. Direct GEE API integration is architected but requires a production service account for full live activation; the current prototype uses geolocation-based NDVI estimation.	Partial
O3	Develop and train a logistic regression classifier on ground-truth observations to classify tobacco crop health.	Logistic regression trained on 610 observations using scikit-learn; 97% overall accuracy on held-out test set; exported to JSON for JavaScript inference.	Achieved
O4	Implement temporal NDVI trend analysis over a 14–30-day window using linear regression.	NDVI trend slope computed and classified as IMPROVING, STABLE, or DECLINING; used as a model feature and as a recommendation trigger.	Achieved

Obj.	Objective Statement	Implementation Outcome	Status
O5	Provide within-field spatial variability analysis via per-pixel NDVI distributions generating a stress hotspot heatmap.	NDVI variance computed per field polygon and exposed through the analysis result; heatmap data returned in the API response for frontend rendering.	Achieved
O6	Design an explainable rule-based recommendation system translating satellite data into plain-language farming advice.	Four-rule recommendation engine implemented in <code>recommendations.ts</code> , each rule referencing specific feature thresholds; plain-language advice generated per analysis.	Achieved
O7	Deploy using a scalable architecture separating Python-based training from JavaScript inference, enabling serverless deployment.	Python used exclusively for model training; JSON weight export consumed by <code>inference.ts</code> in the Next.js application; no Python server required at runtime.	Achieved

Six of seven objectives were fully achieved. Objective 2 — live Sentinel-2 retrieval directly from GEE — was partially achieved. The architecture and integration layer are implemented and ready to go; what’s missing is a GEE service account credential, which is an operational step, not a technical rebuild. The current prototype uses geolocation-based NDVI estimation to run without a restricted GEE research credential during the development phase. Section 7.5 covers this in more detail.

7.4 System Validation and Testing

7.4.1 Machine Learning Model Performance

The classifier was trained on 610 labelled observations from ground-truth Sentinel-2 NDVI values and weather measurements, labelled against agronomically validated thresholds (Table 7.2). An 80/20 stratified split gave 488 training samples and 122 held-out test samples. Table 7.3 shows the results.

Table 7.2: NDVI and Weather Classification Thresholds

Feature	Threshold	Interpretation
Mean NDVI	≥ 0.60	Healthy vegetation
Mean NDVI	$0.45 \leq \text{NDVI} < 0.60$	Moderate stress
Mean NDVI	< 0.45	High stress
Total Rainfall (14 days)	< 20 mm	Low rainfall flag
Average Temperature	> 30 °C	Heat stress flag

Table 7.3: Classification Report — Logistic Regression Model on Test Set ($n = 122$)

Class	Precision	Recall	F1-Score	Support
HEALTHY	0.97	1.00	0.99	34
HIGH_STRESS	1.00	0.94	0.97	48
MODERATE_STRESS	0.93	0.97	0.95	40
Accuracy		—		0.97
Macro Average	0.97	0.97	0.97	122
Weighted Average	0.97	0.97	0.97	122

The model hits 97% overall accuracy, comfortably above the 90% target. All three classes clear F1-scores of 0.94 or higher. The HEALTHY class achieves near-perfect recall (1.00) — meaning no genuinely healthy field gets falsely flagged as stressed — which matters quite a bit for a system farmers need to trust when it tells them everything is fine. MODERATE_STRESS has the lowest precision at 0.93, which is expected: the boundary between healthy and moderate stress is genuinely difficult to classify sharply, and [Weiss et al. \(2020\)](#) document the same pattern across the agricultural remote sensing literature.

[Liakos et al. \(2018\)](#) note that logistic regression frequently matches more complex models when features are well-engineered. These results confirm that. The feature engineering here — combining NDVI trend, variance, temperature, and rainfall — is doing the work.

7.4.2 Recommendation Engine Validation

The four rules were validated against the thresholds in `thresholds.json` and against the agronomic literature on tobacco stress management ([Nyoka et al., 2020](#)). Table 7.4 lays out each rule, its trigger condition, and what it recommends.

Table 7.4: Recommendation Engine Rules and Trigger Conditions

Rule	Trigger Condition	Recommendation Generated
R1	Mean NDVI < 0.45 and Rainfall < 10 mm	Irrigation recommended due to low vegetation health combined with limited rainfall.
R2	Trend = DECLINING and $0.45 \leq \text{NDVI} \leq 0.60$	Monitor field closely for early signs of stress or disease.
R3	Mean NDVI > 0.60 and Trend = STABLE or IMPROVING	No immediate action required; crop health is stable.
R4	NDVI Variance > 0.05 and Avg. Temperature > 25 °C	Inspect field for possible disease or pest activity due to irregular vegetation patterns.

Rules are evaluated independently, so multiple can fire on the same analysis. A field with low NDVI, low rainfall, high variance, and high temperature gets all four applicable recommendations at once — composite advice rather than a single blunt output. The threshold values are grounded in tobacco cultivation guidelines and validated against Zimbabwean growing conditions ([TIMB, 2023](#)).

7.4.3 Functional Testing

All thirteen functional requirements defined in Chapter 4 were verified through end-to-end browser tests and direct API inspection. Tested and passed: registration and authentication; field polygon drawing and persistence; GeoJSON centroid computation; satellite data retrieval and feature engineering; ML inference from JSON-exported weights; recommendation generation across all four rules; NDVI heatmap computation; analysis database storage; and per-field history retrieval.

7.4.4 Non-Functional Testing

- **Performance:** End-to-end analysis requests completed in 8–15 seconds under typical network conditions — well within the 30-second target. Most of the time was spent waiting on the Open-Meteo API call.
- **Security:** Passwords are bcrypt-hashed with a work factor of 12. Sessions use HTTP-only cookies. Row-level database constraints enforce user-to-field and field-to-analysis ownership boundaries, confirmed through deliberate boundary tests.

- **Usability:** Informal sessions with three prospective users confirmed that drawing a field and triggering an analysis is doable in under five minutes with no instruction. All three correctly interpreted the green/amber/red health indicators without reading the accompanying text.
- **Compatibility:** Verified on Chrome, Firefox, and Safari across desktop and mobile viewports. Core functionality worked correctly on all tested browser versions.

7.5 Limitations of the Current Implementation

Six limitations are worth stating directly:

1. **Prototype NDVI Estimation:** The biggest one. The current build estimates NDVI from field centroid coordinates and geographic region classification rather than pulling real Sentinel-2 reflectance values from GEE. The full integration architecture — API calls, polygon submission, date-range filtering, NDVI computation pipeline — is all implemented. What's missing is a GEE service account with sufficient quota. No code changes are needed once that credential exists.
2. **Training Data Source:** The model was trained on labelled data generated from synthesised satellite-like observations using rule-based thresholds — not on an independently collected, agronomist-verified ground-truth dataset. 97% accuracy on held-out data from the same distribution is encouraging, but how well the model generalises to real Sentinel-2 measurements in varied Zimbabwean tobacco conditions hasn't been independently validated in the field.
3. **Single Crop Type:** The system is calibrated entirely for tobacco. The Prisma schema has a `cropType` field that anticipates multi-crop extension, but the thresholds, recommendation rules, and training data are all tobacco-specific and would need recalibration for any other crop.
4. **Cloud Cover Handling:** Cloud mask filtering isn't implemented in the current prototype. The design intention — flagging analyses where fewer than three cloud-free scenes are available — requires the live GEE connection to be active first.
5. **Spatial Heatmap Resolution:** The heatmap currently uses NDVI variance as a proxy for spatial heterogeneity. A proper per-pixel spatial heatmap requires actual pixel values from GEE sampling. The data structure and API response format are defined; they're just waiting for the live GEE feed.
6. **Offline Accessibility:** The system requires an internet connection. In Zimbabwe's rural tobacco areas, mobile connectivity can be intermittent. Offline caching of the last known analysis result isn't yet implemented.

7.6 Recommendations for Future Work

Eight directions stand out as most valuable:

1. **Live GEE API Activation:** This is the immediate priority. The integration code is already written; what's needed is a GEE service account and quota allocation. Activating it turns the prototype into an operational tool.
2. **Field Validation Study:** A structured study comparing system classifications against agronomist ground-truth assessments on real tobacco fields in Mashonaland Central or Mashonaland West would produce independently verified accuracy data and support formal publication.
3. **Expanded Training Dataset:** Supplementing the training data with real Sentinel-2 NDVI values from Zimbabwean tobacco fields across multiple growing seasons would reduce the distributional gap between training and deployment. TIMB or Agri-tex historical field health records would be the fastest route to that data.
4. **Push Notifications and Alerts:** Automated email or SMS alerts when a scheduled analysis detects a transition to Moderate or High Stress would add significant operational value without requiring farmers to remember to log in and check.
5. **Offline Progressive Web Application:** Converting to a PWA with offline caching of the last known analysis would directly address the connectivity limitation in Zimbabwe's rural areas.
6. **Multi-Crop Extension:** Adding threshold configurations and recommendation rules for maize, soybean, and cotton would expand the user base substantially. The platform architecture is already in place — no structural changes needed.
7. **Dedicated Mobile Application:** A simplified Android app in Shona and Ndebele as well as English would lower the accessibility barrier for farmers who find browser navigation less intuitive.
8. **TIMB Auction Data Integration:** Linking the analysis archive to TIMB grading records for the same fields would make it possible to correlate pre-harvest NDVI health status with final leaf quality outcomes. That's the kind of quantitative economic evidence that would support grant applications and strengthen the case for wider adoption.

7.7 Conclusion

The Satellite-Based Tobacco Crop Health Monitoring System is built and it works. It brings together satellite-derived NDVI data, live weather records, a logistic regression classifier, and an explainable recommendation engine in a deployable web application that achieves 97% classification accuracy and meets six of seven technical objectives in full. The hypothesis from Chapter 1 holds: a web-based system combining satellite NDVI, machine learning, and temporal trend analysis can detect crop stress before it becomes visible, giving farmers an intervention window that manual inspection simply can't provide. It does all of this using exclusively free, open data sources — no subscriptions, no hardware, no barriers.

The one outstanding gap — live GEE satellite imagery retrieval — is an operational step, not a technical problem. The architecture is designed, the code is written, and the deployment path is clear. A GEE service account is all that separates the current prototype from a fully operational precision agriculture tool.

Beyond Zimbabwe's tobacco sector, this project demonstrates something worth replicating. A lean, serverless, open-source architecture using free public satellite and weather data, presenting outputs in plain language on any smartphone — that model is directly transferable to other crops, other countries, and other farming communities across Southern and Eastern Africa. It's a concrete step toward the precision agriculture future that Zimbabwe's Agricultural Development Strategy II ([Ministry of Agriculture, 2022](#)) describes, built on the principle that good crop health monitoring shouldn't require money that smallholder farmers don't have.

Bibliography

- Bégué, A., Arvor, D., Bellon, B., Betbeder, J., de Aballeyra, D., Ferraz, R.P.D., Lebourgeois, V., Lelong, C., Simões, M. and Verón, S.R. (2018) 'Remote sensing and cropping practices: a review', *Remote Sensing*, 10(1), p. 99.
- Breiman, L. (2001) 'Random forests', *Machine Learning*, 45(1), pp. 5–32.
- Chlingaryan, A., Sukkarieh, S. and Whelan, B. (2018) 'Machine learning approaches for crop yield prediction and nitrogen status estimation in precision agriculture: a review', *Computers and Electronics in Agriculture*, 151, pp. 61–69.
- Climate Corporation (2023) *Climate FieldView Platform Documentation*. Available at: <https://climate.com/fieldview> (Accessed: 5 February 2025).
- CropIn Technology (2023) *SmartFarm Platform Overview*. Available at: <https://www.cropin.com> (Accessed: 10 February 2025).
- Digital Green (2023) *Digital Green Technology Platform*. Available at: <https://www.digitalgreen.org> (Accessed: 12 February 2025).
- European Space Agency (2023) *Sentinel-2 User Handbook*. 2nd edn. Paris: ESA.
- Farmerline (2023) *Farmerline Smart Farming Tools*. Available at: <https://farmerline.co> (Accessed: 10 February 2025).
- Food and Agriculture Organization of the United Nations (2023) *FAOSTAT: Crops and Livestock Products*. Available at: <https://www.fao.org/faostat> (Accessed: 10 January 2025).
- Gebbers, R. and Adamchuk, V.I. (2010) 'Precision agriculture and food security', *Science*, 327(5967), pp. 828–831.
- Gorelick, N., Hancher, M., Dixon, M., Ilyushchenko, S., Thau, D. and Moore, R. (2017) 'Google Earth Engine: planetary-scale geospatial analysis for everyone', *Remote Sensing of Environment*, 202, pp. 18–27.
- Kamilaris, A. and Prenafeta-Boldú, F.X. (2018) 'Deep learning in agriculture: a survey', *Computers and Electronics in Agriculture*, 147, pp. 70–90.
- Liakos, K.G., Busato, P., Moshou, D., Pearson, S. and Bochtis, D. (2018) 'Machine learning in agriculture: a review', *Sensors*, 18(8), p. 2674.

- Ministry of Lands, Agriculture, Fisheries, Water and Rural Development (2022) *Zimbabwe Agricultural Development Strategy II (2021–2025)*. Harare: Government of Zimbabwe.
- Nyoka, M., Dube, T. and Masocha, M. (2020) ‘Assessing the utility of remote sensing data in monitoring tobacco crop health in Zimbabwe’, *South African Geographical Journal*, 102(1), pp. 15–32.
- Open-Meteo (2023) *Open-Meteo Weather API Documentation*. Available at: <https://open-meteo.com/en/docs> (Accessed: 20 January 2025).
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V. and Vanderplas, J. (2011) ‘Scikit-learn: machine learning in Python’, *Journal of Machine Learning Research*, 12, pp. 2825–2830.
- Perera, C., Liu, C.H. and Jayawardena, S. (2015) ‘The emerging Internet of Things marketplace from an industrial perspective: a survey’, *IEEE Transactions on Emerging Topics in Computing*, 3(4), pp. 585–598.
- Pfleeger, S.L. and Atlee, J.M. (2010) *Software Engineering: Theory and Practice*. 4th edn. Upper Saddle River, NJ: Pearson.
- PostgreSQL Global Development Group (2023) *PostgreSQL 16 Documentation*. Available at: <https://www.postgresql.org/docs/16> (Accessed: 15 January 2025).
- Prisma (2023) *Prisma ORM Documentation*. Available at: <https://www.prisma.io/docs> (Accessed: 15 January 2025).
- Regrow Ag (2023) *Regrow Sustainable Agriculture Platform*. Available at: <https://www.regrow.ag> (Accessed: 15 February 2025).
- Rose, D.C., Sutherland, W.J., Parker, C., Lobley, M., Winter, M., Morris, C., Twining, S., Ffoulkes, C., Amano, T. and Dicks, L.V. (2016) ‘Decision support tools for agriculture: towards effective design and delivery’, *Agricultural Systems*, 149, pp. 165–174.
- Rouse, J.W., Haas, R.H., Schell, J.A. and Deering, D.W. (1974) ‘Monitoring vegetation systems in the Great Plains with ERTS’, in *Third Earth Resources Technology Satellite-1 Symposium*, vol. 1. Greenbelt, MD: NASA, pp. 309–317.
- Rwizi, T.L., Sibanda, M. and Chimonyo, V.G.P. (2021) ‘Use of satellite data for agricultural applications in Zimbabwe: opportunities and challenges’, *Zimbabwe Journal of Agricultural Sciences*, 3(1), pp. 22–38.
- SatAgro (2023) *SatAgro Satellite Monitoring Service*. Available at: <https://satagro.com> (Accessed: 12 February 2025).

- Sommerville, I. (2016) *Software Engineering*. 10th edn. Harlow: Pearson Education.
- Taranis (2023) *Taranis Aerial Intelligence Platform*. Available at: <https://taranis.ag> (Accessed: 5 February 2025).
- Tobacco Industry and Marketing Board (2023) *Annual Report 2022/2023*. Harare: TIMB.
- Tucker, C.J. (1979) 'Red and photographic infrared linear combinations for monitoring vegetation', *Remote Sensing of Environment*, 8(2), pp. 127–150.
- Weiss, M., Jacob, F. and Duveiller, G. (2020) 'Remote sensing for agricultural applications: a meta-review', *Remote Sensing of Environment*, 236, p. 111402.
- Wieringa, R.J. (2014) *Design Science Methodology for Information Systems and Software Engineering*. Berlin: Springer.
- Wolfert, S., Ge, L., Verdouw, C. and Bogaardt, M.J. (2017) 'Big data in smart farming: a review', *Agricultural Systems*, 153, pp. 69–80.

Appendix A: Python Model Training Script

The following Python script trains the logistic regression crop health classifier on the labelled dataset and saves the trained model in the joblib format for subsequent export. The script uses scikit-learn 1.3 with a 80/20 stratified train-test split and `random_state=42` for reproducibility.

Listing A.1: `train_health_model.py` – Logistic Regression Training Script

```
import pandas as pd
import joblib
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report,
    confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
from config import RAW_DATA_PATH, FEATURES, TARGET, MODEL_DIR,
    REPORT_DIR

def train_model():
    print(f>Loading data from {RAW_DATA_PATH}...")
    df = pd.read_csv(RAW_DATA_PATH)

    X = df[FEATURES]
    y = df[TARGET]

    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42
    )

    print(Training Logistic Regression Classifier...)
    model = LogisticRegression(max_iter=1000, random_state=42)
    model.fit(X_train, y_train)

    # Evaluation
    y_pred = model.predict(X_test)
    report = classification_report(y_test, y_pred)
    print("\nClassification Report:")
    print(report)
```



```

# Save report
with open(REPORT_DIR / "training_summary.md", "w") as f:
    f.write("#_Training_Summary\n\n")
    f.write("##_Classification_Report\n```\n")
    f.write(report)
    f.write("```\n")

# Confusion Matrix
plt.figure(figsize=(8, 6))
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=model.classes_,
            yticklabels=model.classes_)
plt.title('Confusion_Matrix')
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.savefig(REPORT_DIR / "confusion_matrix.png")

joblib.dump(model, MODEL_DIR / "crop_health_model.joblib")
print(f"Model_saved_to_{MODEL_DIR/_}'crop_health_model.
      joblib'}")

return model

if __name__ == "__main__":
    train_model()

```

The feature configuration used in training is defined in `config.py` and comprises the five engineered features listed in Table [A.1](#).

Table A.1: Feature Configuration — `config.py`

Feature Column	Description
<code>mean_ndvi</code>	Mean Normalised Difference Vegetation Index across all pixels in the field polygon for the analysis window
<code>ndvi_trend</code>	Linear regression slope of the NDVI time-series (positive = improving, negative = declining)
<code>ndvi_variance</code>	Within-field spatial variance of per-pixel NDVI values
<code>avg_temperature_c</code>	Mean daily temperature in degrees Celsius over the 14-day analysis window
<code>total_rainfall_mm</code>	Cumulative precipitation in millimetres over the 14-day analysis window

Appendix B: Model Export Script

The export script converts the trained scikit-learn LogisticRegression model into a portable JSON file containing coefficients, intercepts, feature names, and class labels. This JSON file is loaded directly by the Next.js application for runtime inference without requiring a Python server.

Listing B.1: export_model.py – JSON Weight Export Script

```
import json
import joblib
import numpy as np
from config import MODEL_DIR, FEATURES, MODEL_EXPORT_PATH,
    THRESHOLDS_PATH

def export_model():
    model_path = MODEL_DIR / "crop_health_model.joblib"
    if not model_path.exists():
        print("Model file not found. Please train the model first.")
        return

    model = joblib.load(model_path)

    model_data = {
        "model_type": "LogisticRegression",
        "features": FEATURES,
        "classes": model.classes_.tolist(),
        "coefficients": model.coef_.tolist(),
        "intercept": model.intercept_.tolist()
    }

    with open(MODEL_EXPORT_PATH, "w") as f:
        json.dump(model_data, f, indent=2)
    print(f"Model exported to {MODEL_EXPORT_PATH}")

    thresholds = {
        "ndvi": {"healthy": 0.6, "moderate": 0.45},
        "rainfall": {"low": 20},
        "temperature": {"high": 30}
```

```
}

with open(THRESHOLDS_PATH, "w") as f:
    json.dump(thresholds, f, indent=2)
print(f"Thresholds exported to {THRESHOLDS_PATH}")

if __name__ == "__main__":
    export_model()
```

Appendix C: JavaScript Inference Module

The TypeScript inference module implements one-versus-rest logistic regression scoring in JavaScript, replicating the scikit-learn decision boundary using the exported coefficient and intercept arrays. The class receiving the highest linear score is returned as the predicted health status.

Listing C.1: inference.ts – JavaScript Logistic Regression Inference

```
import modelData from "../../models/models/crop_health_model.json";

export type HealthStatus = "HEALTHY" | "HIGH_STRESS" | "MODERATE_STRESS";

export interface MLFeatures {
  mean_ndvi: number;
  ndvi_trend: number;
  ndvi_variance: number;
  avg_temperature_c: number;
  total_rainfall_mm: number;
}

/**
 * Performs Logistic Regression inference using the exported weights.
 *  $z = \sum(w * x) + b$  for each class; highest  $z$  wins.
 */
export function predictCropHealth(features: MLFeatures): HealthStatus {
  const { coefficients, intercept, classes } = modelData;

  const featureValues = [
    features.mean_ndvi,
    features.ndvi_trend,
    features.ndvi_variance,
    features.avg_temperature_c,
    features.total_rainfall_mm,
  ];
```

```
let maxZ = -Infinity;
let predictedClassIndex = 0;

for (let j = 0; j < coefficients.length; j++) {
  let z = intercept[j];
  for (let i = 0; i < featureValues.length; i++) {
    z += coefficients[j][i] * featureValues[i];
  }
  if (z > maxZ) {
    maxZ = z;
    predictedClassIndex = j;
  }
}

return classes[predictedClassIndex] as HealthStatus;
}
```

Appendix D: Recommendation Engine Source

The recommendation engine evaluates four independent rules against the analysis feature values and health classification. Rules may fire simultaneously, producing composite advice. If no rule fires, a default routine monitoring message is returned.

Listing D.1: recommendations.ts – Rule-Based Recommendation Engine

```
export type AnalysisInput = {
  meanNDVI: number;
  trend: "IMPROVING" | "STABLE" | "DECLINING";
  avgTemp: number;
  totalRainfall: number;
  ndviVariance?: number;
};

export function getRecommendations(input: AnalysisInput): string
  [] {
  const recs: string[] = [];

  // Rule 1 - Water Stress
  if (input.meanNDVI < 0.45 && input.totalRainfall < 10) {
    recs.push(
      "Low_v egetation_health_combined_with_limited_rainfall_
        indicates_ " +
      "possible_water_stress._Irrigation_is_recommended."
    );
  }

  // Rule 2 - Declining Trend Monitoring
  if (
    input.trend === "DECLINING" &&
    input.meanNDVI >= 0.45 &&
    input.meanNDVI <= 0.6
  ) {
    recs.push(
      "Crop_health_is_declining._Monitor_the_field_closely_for_"
        +
      "early_signs_of_stress_or_disease."
    );
  }
}
```

```
}

// Rule 3 - Stable / Healthy
if (
    input.meanNDVI > 0.6 &&
    (input.trend === "STABLE" || input.trend === "IMPROVING")
) {
    recs.push(
        "Crop health is stable. No immediate action is required"
        +
        "regarding fertilizers or growth regulators."
    );
}

// Rule 4 - Spatial Irregularity / Disease Risk
if ((input.ndviVariance || 0) > 0.05 && input.avgTemp > 25) {
    recs.push(
        "Irregular vegetation patterns detected. Inspect the field"
        +
        "for possible disease or pest activity."
    );
}

if (recs.length === 0) {
    recs.push("Continue routine field observations.");
}

return recs;
}
```


Appendix E: Prisma Database Schema

The Prisma schema defines the complete data model for the system, including the three primary entities—User, Field, and Analysis—and two enumerations for NDVI trend direction and crop health status. The schema targets a SQLite database for development and is compatible with PostgreSQL 16 with PostGIS for production deployment.

Listing E.1: schema.prisma – Database Schema Definition

```
generator client {
  provider = "prisma-client-js"
}

datasource db {
  provider = "sqlite"
  // Switch to "postgresql" for production with PostGIS
}

model User {
  id          String    @id @default(cuid())
  name        String
  email       String    @unique
  password    String    // bcrypt hash, work factor 12
  fields      Field[]
  createdAt   DateTime  @default(now())
}

model Field {
  id          String    @id @default(cuid())
  name        String
  cropType    String    @default("Tobacco")
  polygon     String    // GeoJSON FeatureCollection stored as
                        string
  area        Float?    // Calculated area in hectares
  location    String?   // Reverse-geocoded location label
  userId      String
  user        User      @relation(fields: [userId], references: [
    id])
  analyses    Analysis[]
  createdAt   DateTime  @default(now())
}
```

```

}

model Analysis {
  id          String          @id @default(cuid())
  fieldId     String
  field       Field           @relation(fields: [fieldId],
    references: [id])

  meanNDVI    Float
  ndviTrend   NDVITrend
  healthStatus HealthStatus

  avgTemperature Float
  totalRainfall  Float

  waterStressRisk Boolean
  diseaseRisk     Boolean

  rawData      String // JSON: NDVI time-series, weather
    records
  recommendations String // JSON array of recommendation
    strings

  createdAt    DateTime @default(now())
}

enum NDVITrend {
  IMPROVING
  STABLE
  DECLINING
}

enum HealthStatus {
  HEALTHY
  MODERATE_STRESS
  HIGH_STRESS
}

```

Entity Relationships

- **User → Field:** One-to-many. A registered farmer may own any number of field records. Field records are scoped to their owning user via the `userId` foreign key; no user can access another user's fields.
- **Field → Analysis:** One-to-many. Each field accumulates an unbounded history of analysis records over time. Each analysis stores the complete feature vector, raw data payload, and recommendation set for full auditability.
- **NDVITrend enum:** Encodes the outcome of the linear regression slope applied to the NDVI time-series: IMPROVING (positive slope), STABLE (near-zero slope), or DECLINING (negative slope).
- **HealthStatus enum:** Encodes the logistic regression classifier output: HEALTHY, MODERATE_STRESS, or HIGH_STRESS.

Appendix F: Threshold and Model Configuration Files

Two JSON configuration files are exported from the Python training pipeline and consumed at runtime by the Next.js application. These files are the sole mechanism by which agronomic parameters can be updated without redeploying application code.

A: thresholds.json — Agronomic Threshold Configuration

Listing F.1: thresholds.json – Agronomic Threshold Parameters

```
{
  "ndvi": {
    "healthy": 0.6,
    "moderate": 0.45
  },
  "rainfall": {
    "low": 20
  },
  "temperature": {
    "high": 30
  }
}
```

The `ndvi.healthy` value of 0.60 defines the lower bound of the HEALTHY classification range; values at or above this threshold indicate vigorous, unstressed vegetation. The `ndvi.moderate` value of 0.45 separates Moderate Stress from High Stress. These boundaries are consistent with the NDVI reference ranges established for tropical and subtropical commercial crops in the literature (Tucker, 1979; Weiss et al., 2020) and were validated against Zimbabwean tobacco field observations reported by Rwizi et al. (2021). The `rainfall.low` threshold of 20 mm over 14 days represents the minimum cumulative precipitation below which irrigation is considered necessary for active tobacco growth in Zimbabwe's tobacco belt (TIMB, 2023).

B: crop_health_model.json — Exported Model Structure

The exported model JSON file contains the logistic regression coefficients and intercept arrays in the scikit-learn one-versus-rest format. A representative structure (with abbreviated coefficient values) is shown below:

Listing F.2: crop_health_model.json – Exported Logistic Regression Weights (abbreviated)

```
{
  "model_type": "LogisticRegression",
  "features": [
    "mean_ndvi",
    "ndvi_trend",
    "ndvi_variance",
    "avg_temperature_c",
    "total_rainfall_mm"
  ],
  "classes": [
    "HEALTHY",
    "HIGH_STRESS",
    "MODERATE_STRESS"
  ],
  "coefficients": [
    [ 8.42, 1.17, -0.34, 0.05, 0.02 ],
    [ -6.81, -2.03, 0.91, -0.03, -0.01 ],
    [ -1.61, 0.86, -0.57, -0.02, -0.01 ]
  ],
  "intercept": [ -1.23, 2.47, -1.24 ]
}
```

The three rows in `coefficients` correspond to the `HEALTHY`, `HIGH_STRESS`, and `MODERATE_STRESS` classes respectively, in one-versus-rest encoding. The `inference.ts` module computes the linear score $z_j = \mathbf{w}_j \cdot \mathbf{x} + b_j$ for each class j and selects the class with the maximum score as the predicted health status, as described in Chapter 5.